

# Beyond Prepare-and-Measure: The Cost of Trustless Detectors in Quantum Key Distribution

Linus Kelsey

Aug 2026

## Abstract

Quantum Key Distribution (QKD) was first put forward in 1984 by Bennett and Brassard through their foundational proposal for bipartite secure key exchange taking advantage of quantum mechanics, now commonly known as BB84. In the following decades, a number of inspired methods have emerged to guarantee shared secrecy. In particular, Measurement-Device-Independent (MDI) QKD, presented in 2012, improves BB84 security assumptions by eliminating the need to trust photon detectors. In order to be able to deploy QKD at scale, providers must map the trade-offs of either protocol, and comparisons of the two in secure information exchange rate (key rate), cost and scalability allow for efficient application of this maturing technology. In this report, we aim to contribute to the solution of this problem through use of NetSquid, a quantum network discrete event simulation (DES) engine, and an algebraic cost model, to simulate and compare network behaviour for both MDI-QKD and BB84. We see that the price to pay for MDI's improved security is a drop in key rate, of up to 10 times, in modest metropolitan networks under realistic hardware assumptions.

# Contents

|          |                                               |           |
|----------|-----------------------------------------------|-----------|
| <b>1</b> | <b>Introduction</b>                           | <b>4</b>  |
| <b>2</b> | <b>Shared Quantum Secrecy</b>                 | <b>6</b>  |
| 2.1      | BB84 . . . . .                                | 6         |
| 2.2      | Measurement-Device-Independent QKD . . . . .  | 7         |
| 2.3      | Theoretical Performance Bounds . . . . .      | 9         |
| 2.4      | Network Topologies . . . . .                  | 9         |
| <b>3</b> | <b>Methodology</b>                            | <b>11</b> |
| 3.1      | Performance Analysis . . . . .                | 11        |
| 3.1.1    | NetSquid Simulation Engine . . . . .          | 12        |
| 3.1.2    | Hardware Model . . . . .                      | 12        |
| 3.1.3    | Bipartite Single-Link Analysis . . . . .      | 15        |
| 3.1.4    | Multi-User Network Analysis . . . . .         | 17        |
| 3.2      | Cost Model . . . . .                          | 20        |
| 3.2.1    | Component Counts . . . . .                    | 20        |
| 3.2.2    | Symbolic Cost Formulae . . . . .              | 21        |
| 3.3      | Scalability Analysis . . . . .                | 22        |
| <b>4</b> | <b>Results</b>                                | <b>24</b> |
| 4.1      | Point-to-Point Performance . . . . .          | 24        |
| 4.2      | Network Performance . . . . .                 | 31        |
| 4.3      | Cost . . . . .                                | 31        |
| 4.4      | Deployment Recommendations . . . . .          | 31        |
| <b>5</b> | <b>Discussion</b>                             | <b>32</b> |
| 5.1      | The Security–Rate Trade-off . . . . .         | 32        |
| 5.2      | Cost Model Limitations . . . . .              | 33        |
| 5.3      | Simulation Model Limitations . . . . .        | 34        |
| 5.4      | Implications for Protocol Selection . . . . . | 37        |
| <b>6</b> | <b>Conclusion</b>                             | <b>38</b> |
| 6.1      | Summary of Contributions . . . . .            | 38        |
| 6.2      | Findings . . . . .                            | 38        |
| 6.3      | Future Work . . . . .                         | 39        |

|          |                                                          |           |
|----------|----------------------------------------------------------|-----------|
| <b>A</b> | <b>Additional Point-to-Point Parameter Sweeps</b>        | <b>41</b> |
| <b>B</b> | <b>Relay Placement Optimisation</b>                      | <b>43</b> |
| B.1      | Derivation of the Coupled Update . . . . .               | 43        |
| B.2      | Worked Example: Centroid vs. Geometric Median . . . . .  | 44        |
| B.3      | Geometric Comparison Across Network Topologies . . . . . | 45        |
| B.4      | Network-Level Comparison Under Organic Growth . . . . .  | 47        |
| B.5      | Open Questions . . . . .                                 | 47        |

# 1 Introduction

The advent of large-scale quantum computing poses a fundamental threat to classical public-key cryptography. Modern key-exchange standards, including RSA [1] and Elliptic Curve Cryptography [2, 3], derive their security from the computational hardness of integer factorisation and the discrete logarithm problem. Shor’s algorithm [4], published in 1997, demonstrated that both problems can be solved in polynomial time on a sufficiently powerful quantum computer, rendering these standards vulnerable to a quantum-capable adversary. Post-Quantum Cryptography (PQC) responds to this threat by identifying classically hard problems not susceptible to known quantum speed-ups; Quantum Key Distribution (QKD) takes the complementary route, using quantum mechanics to guarantee shared secrecy rather than computational assumption.

QKD was first conceived by Wiesner [5] in the 1970s and given its first operational form by Bennett and Brassard in 1984 [6]. In the decades since, QKD has matured into one of the most field-tested quantum technologies, with national-scale fibre deployments serving over 150 users across thousands of kilometres [7] and metropolitan testbeds demonstrating sustained multi-vendor operation [8, 9]. This trajectory has motivated a growing body of research into how QKD protocols should be chosen and deployed at scale. This work focuses on two protocols that represent a natural axis of comparison: BB84, the original Prepare-and-Measure scheme, and Measurement-Device-Independent QKD (MDI-QKD), proposed by Lo, Curty and Qi in 2012 [10]. Their protocols and theoretical performance bounds are introduced in Section 2.

MDI-QKD improves upon BB84’s security model by removing the need to trust photon detectors, an assumption that is difficult to verify in practice and has been exploited in hardware-level attacks. This advantage, however, comes at a real cost in key rate: at metropolitan distances, BB84 consistently delivers rates an order of magnitude higher than MDI under comparable hardware conditions [11, 12]. For network providers selecting between the two protocols, this trade-off must be mapped carefully across distance, topology, and network scale. The question is not simply which protocol is more secure, but under what deployment conditions MDI’s security advantage justifies its rate penalty, and at what economic and energetic cost.

Despite the commercial momentum behind both protocols — evidenced by a progression of field deployments from early national testbeds [7–9, 13, 14] through metropolitan and long-haul fibre networks [11, 15] to dedicated MDI-QKD systems [16–18], heterogeneous multi-vendor network trials [19], and European infrastructure programmes [20] — the quantitative basis for choosing between BB84 and MDI-QKD for a given application

remains poorly characterised. Simulation of MDI-QKD at metropolitan scale has not systematically explored the interplay between topology, relay placement, and hardware parameters across a sufficiently wide parameter space to support confident protocol selection [21]. Cost models for MDI networks are largely absent from the literature: whilst integer linear programming approaches have been developed for PM and entanglement-based QKD infrastructure [22], their extension to MDI has not been carried out. Energetic footprint, an increasingly central criterion in telecommunications infrastructure procurement, has only recently begun to receive attention in the QKD context [23]. A rigorous parametric comparison across key rate, cost, and scalability is absent.

This work addresses these gaps through a two-part methodology. First, we implement and simulate both protocols within NetSquid [24], a discrete-event quantum network simulation engine, across a ten-layer hardware model spanning ideal to physically realistic conditions. Second, we derive deployment cost formulae algebraically using symbolic per-component parameters—cost per photon source  $c_s$ , per single-photon detector  $c_d$ , and per km of installed dark fibre  $c_f$ —expressing each protocol’s total infrastructure cost as a closed-form function of user count and topology. Since manufacturer prices for quantum dot sources and single-photon detectors are not publicly available, absolute CapEx comparison is deferred to the parameter space; the structural scaling difference between protocols is the primary analytical result. We compare BB84 and MDI-QKD across three dimensions: key rate as a function of distance and hardware parameters; relative deployment cost as a function of network size; and scalability, defined as the marginal key rate when adding users to a fixed metropolitan topology. The remainder of this report is structured as follows. Section 2 introduces the protocols, their hardware underpinnings, and the theoretical performance bounds that govern them. Section 3 describes the simulation architecture and cost model. Section 4 presents results across all three comparison axes, and Section 6 concludes with deployment recommendations and directions for future work.

## 2 Shared Quantum Secrecy

Quantum Key Distribution refers to any communications protocol in which two or more parties use quantum channels, the sharing of quantum states, together with authenticated classical channels and post-processing, to establish a shared secret cryptographic key. The security of such a scheme is guaranteed by the laws of quantum mechanics, providing information-theoretic security against any eavesdropper whose capabilities are bounded by quantum physics. Since Bennett and Brassard’s foundational paper in 1984 [6], numerous variants of QKD have been proposed. This work focuses on two: the original Prepare-and-Measure scheme (BB84) and Measurement-Device-Independent QKD (MDI-QKD), proposed by Lo, Curty and Qi in 2012 [10].

### 2.1 BB84

BB84 is representative of the class of Prepare-and-Measure (PM) protocols. Alice begins by preparing a string of  $n$  photons, each encoded in one of four polarisation states ( $0^\circ$ ,  $45^\circ$ ,  $90^\circ$  or  $135^\circ$ ) selected uniformly at random. In the qubit formalism, these correspond to  $|0\rangle$ ,  $|+\rangle = \frac{1}{\sqrt{2}}(|0\rangle + |1\rangle)$ ,  $|1\rangle$  and  $|-\rangle = \frac{1}{\sqrt{2}}(|0\rangle - |1\rangle)$ , respectively. Alice selects a basis ( $\{|0\rangle, |1\rangle\}$  or  $\{|+\rangle, |-\rangle\}$ ) for each qubit uniformly and independently, then uniformly and independently again, she applies a bit flip within that basis before transmission.

She transmits these quantum states to Bob over a quantum channel. Bob measures each photon in a uniformly and independently random basis. Alice and Bob then use an authenticated classical channel to compare basis choices and discard all photons measured in mismatching bases, producing a raw, sifted key. A portion of this sifted key is compared over the classical channel to estimate the Quantum Bit Error Rate (QBER): the fraction of photon pairs measured in the same basis that nevertheless disagree. A QBER below 11% is taken as evidence that the channel has not been compromised [25]. Alice and Bob then apply error correction and privacy amplification (hash-based post-processing that shortens the raw key, removing any partial information an eavesdropper may have gathered) to produce the final shared secret key. Figure 1 illustrates the basic bipartite logical layer.




Figure 1: Communication diagram for point-to-point BB84. Alice sends quantum states to Bob; basis reconciliation and post-processing occur over authenticated classical channels.

Security proofs for BB84, and QKD generally, are conditional on the modelling assumptions made about devices, network, and attacker capabilities. A particular vulnerability in practice arises from weak coherent pulse (WCP) laser sources, whose photon number per pulse follows a Poisson distribution. Multi-photon pulses allow a Photon Number Splitting (PNS) attack [26], in which Eve intercepts and holds trailing photons until bases are announced, extracting key information without disturbing the QBER. The decoy-state technique [27] mitigates this by varying pulse intensity randomly, enabling Alice and Bob to tightly bound the single-photon contribution to their key.

## 2.2 Measurement-Device-Independent QKD

MDI-QKD addresses a fundamental vulnerability shared by all PM schemes: the need to trust the measurement devices at Bob’s node. Side-channel attacks on detectors represent a significant practical threat, as the gap between idealised device models assumed in security proofs and real hardware is difficult to close. MDI removes detectors entirely from the trust boundary by introducing an untrusted third-party relay, Charlie, who performs all measurement operations.

In MDI-QKD, both Alice and Bob independently prepare photon strings in the manner of BB84 and transmit these to Charlie. Charlie performs a Bell State Measurement (BSM) on each arriving photon pair, projecting onto the two-qubit Bell basis:

$$|\phi^\pm\rangle = \frac{|00\rangle \pm |11\rangle}{\sqrt{2}}, \quad |\psi^\pm\rangle = \frac{|01\rangle \pm |10\rangle}{\sqrt{2}}.$$

Charlie announces the measurement outcome over a classical channel to both Alice and Bob. Since the BSM measures only the *joint correlation* between the two incoming photons rather than their individual polarisations, an eavesdropper with full control of Charlie’s apparatus can learn nothing about Alice’s or Bob’s individual bit values. The physical mechanism underlying the BSM is Hong-Ou-Mandel (HOM) interference [28]: two indistinguishable photons arriving simultaneously at a 50:50 beam splitter always exit the same output port due to bosonic quantum interference. Coincident clicks at the downstream po-

larising beam splitters then identify the projected Bell state. Photon indistinguishability is therefore a prerequisite for MDI-QKD: timing jitter, spectral mismatch, and polarisation misalignment all reduce HOM visibility and degrade BSM success probability. The bipartite logical layer for MDI-QKD is shown in Figure 2.



Figure 2: Communication diagram for point-to-point MDI-QKD. Both Alice and Bob transmit quantum states to Charlie; Charlie performs a BSM and announces the result classically to both parties.

Alice and Bob sift by matching bases as in BB84. Upon receiving the BSM outcome, each party knows their own basis choice and bit value, and can deduce the other’s bit via the outcome. Depending on the Bell state announced, one party selectively bit-flips to ensure agreement, as summarised in Table 1. In the linear-optical BSM implementation, only two of the four Bell states ( $|\psi^+\rangle$  and  $|\psi^-\rangle$ ) are unambiguously distinguishable [10]; detection events corresponding to  $|\phi^\pm\rangle$  are discarded as inconclusive. Privacy amplification and error correction then proceed as in BB84.

| BSM outcome       | $ \psi^+\rangle$ | $ \psi^-\rangle$ |
|-------------------|------------------|------------------|
| Diagonal basis    | No bit flip      | Bit flip         |
| Rectilinear basis | Bit flip         | Bit flip         |

Table 1: Bit-flip rules upon receipt of BSM outcome [10]. Alice and Bob agree in advance which party applies the flip; if both flip, there is no net change.

The abstract operation of the BSM in quantum information terms is shown in Figure 3: a CNOT gate followed by a Hadamard and measurement. The physical implementation uses a 50:50 beam-splitter and polarising beam-splitters, with the Hong-Ou-Mandel effect ensuring that only coincidence detection events, two simultaneous clicks at the detector array, constitute valid BSM outcomes.




Figure 3: Bell State Measurement circuit: CNOT gate followed by Hadamard, then measurement in the computational basis.



MDI-QKD’s security advantage comes at a cost. Key generation requires a successful coincidence detection at Charlie – both photons from Alice and Bob must arrive and be detected within the same time window. This requirement, combined with independent connection losses on each arm, imposes a rate penalty relative to BB84’s single-link loss [10]. MDI’s QBER tolerance is also more sensitive than BB84’s, as two independent quantum channels each contribute noise, enlarging the attack surface accounted for by security proofs.

## 2.3 Theoretical Performance Bounds

The maximum achievable secure key rate in any point-to-point (no quantum repeater) QKD scheme is bounded by the Pirandola–Laurenza–Ottaviani–Banchi (PLOB) bound [29]. This fundamental result establishes that, against all attacks permitted by quantum physics, the key rate of a repeater-less link scales at most linearly with channel transmittance  $\eta$ :

$$R \leq -\log_2(1 - \eta) \approx \eta \log_2 e \quad \text{for } \eta \ll 1.$$

Since transmittance decays exponentially with fibre length  $L$  as  $\eta = 10^{-\alpha L/10}$  (with  $\alpha$  the fibre attenuation coefficient in dB/km), the PLOB bound implies that key rates must fall exponentially with distance. Both BB84 and standard MDI-QKD are subject to this bound.

Twin-Field QKD (TF-QKD), proposed in 2018 by Lucamarini et al. [30], exploits the same untrusted-relay tenet as MDI but replaces the two-photon BSM with single-photon interference. The result is that key rates scale as  $\mathcal{O}(\sqrt{\eta})$  rather than  $\mathcal{O}(\eta)$ , exceeding the PLOB bound and extending the distance at which secure communication is achievable. TF-QKD thus represents the most practically significant near-term stepping stone beyond the repeater-less limit, and motivates the question central to this work, of precisely where and under what conditions MDI-QKD’s security advantage over BB84 justifies its rate penalty, before that regime is superseded by TF-QKD or, eventually, quantum repeaters.

## 2.4 Network Topologies

When QKD is deployed beyond a single point-to-point link, network topology becomes a primary driver of cost and scalability. BB84 requires a dedicated quantum channel between every pair of users: for  $N$  users, a full mesh demands  $\binom{N}{2} = \mathcal{O}(N^2)$  fibre spans, each carrying an independent quantum signal. MDI-QKD’s architecture naturally admits a star topology instead: all users direct their pulses to a shared measurement node (Charlie), which performs the Bell-state measurement without accessing the key material. Crucially,

Charlie need not be trusted—the measurement-device-independence property is preserved at network scale, so the relay introduces no new security assumption. For  $K$  relay nodes each serving a cluster of users, the MDI network requires  $\mathcal{O}(N)$  user–relay fibre spans and  $\mathcal{O}(K^2)$  inter-relay backbone links, both substantially below the  $\mathcal{O}(N^2)$  full mesh for  $K \ll N$ . This structural asymmetry—summarised in the component-count model of Section 3.2—is the central economic and scalability argument evaluated in this work.

### 3 Methodology

We compare BB84 and MDI-QKD across three dimensions: performance, deployment cost, and scalability. Together these axes capture the trade-off a network provider faces when selecting between the two protocols for a given metropolitan deployment.

*Performance.* The primary metric is secure key rate, measured in bits per second: the length of the final distilled key divided by the time taken to produce it. Since all QKD-generated keys are information-theoretically equivalent in security regardless of protocol, key rate is the sole performance discriminant between BB84 and MDI-QKD. We simulate key rate as a function of distance and individual hardware parameters using NetSquid [24], a discrete-event quantum network simulation engine, across a ten-layer hardware model that progressively introduces realistic physical imperfections from an ideal baseline.

*Cost.* Deployment cost is modelled algebraically using symbolic per-component parameters:  $c_s$  (photon source),  $c_d$  (single-photon detector), and  $c_f$  (fibre per km). The number of each component is determined by the network topology under each protocol. Since manufacturer prices for quantum dot sources and single-photon detectors are not publicly available at the precision required for quantitative comparison, we express total cost symbolically and focus on the *relative* cost structure between protocols. This reveals a structural crossover: BB84 fibre cost scales as  $\mathcal{O}(N^2)$  whilst MDI scales as  $\mathcal{O}(N)$  at fixed relay count  $K$ , so MDI becomes cost-competitive above a threshold user count  $N^*$  whose value depends on the ratio  $c_d/c_f$  and network geometry.

*Scalability.* We define scalability as the marginal key rate incurred when adding a user to an existing network. Since BB84 requires a direct link between every user pair whilst MDI routes traffic through relay nodes, the two protocols exhibit fundamentally different scaling behaviour. We quantify this by sweeping user count at fixed topology parameters and extracting per-user marginal key rates.

#### 3.1 Performance Analysis

The primary performance metric is the secure key rate  $R$ , measured in bits per second: the key length extracted by a simulation run divided by the simulation duration. Since all QKD-generated keys are information-theoretically equivalent in security regardless of protocol, key rate is the sole performance discriminant between BB84 and MDI-QKD for a given network. A symmetric QBER threshold of 11% is applied: simulation runs with QBER above this value are discarded as failures [25]. All key rates are computed in the asymptotic regime, without finite-key corrections; this overestimates extractable key at

low photon count rates but is standard in comparative simulation studies at metropolitan scale.

### 3.1.1 NetSquid Simulation Engine

NetSquid [24] is a discrete-event simulation (DES) engine for quantum networks developed at QuTech, chosen for this work for three reasons. First, NetSquid natively tracks qubit states as density matrices and exposes standard gate operations (`ns.X`, `ns.H`, `ns.CNOT`) and projective measurement (`ns.qubits.measure`), enabling BB84 encoding and the MDI Bell state measurement to be implemented directly as quantum operations rather than hand-coded probability tables. Second, the DES clock provides per-trial simulation timestamps, making key rate computable from the elapsed simulation time between photon emission and protocol completion; propagation delay is introduced naturally by the channel length parameters of `QuantumChannel` and `ClassicalChannel` objects via a `HybridDelayModel`. Third, the `NodeProtocol` abstraction models each party’s logic as an event-driven coroutine that yields at `await_port_input` calls until a message arrives; this allows Alice, Bob, and Charlie’s asynchronous exchanges (photon emission, index tagging, BSM announcement, basis sifting) to be written sequentially without explicit thread management, and maps directly onto the protocol descriptions in Section 2.

NetSquid organises a simulation into three layers of abstraction. *Components* represent physical hardware: `QuantumChannel` objects carry photons between nodes, applying noise models (loss, dephasing, delay) as Kraus operators; `ClassicalChannel` objects carry classical messages for basis reconciliation and BSM announcements. *Nodes* represent network parties (Alice, Bob, Charlie), each holding named ports through which components connect; nodes own their local quantum processors and memory registers. *Protocols* are event-driven coroutines that run on nodes, yielding control to the DES scheduler at each `await` expression until a subscribed event fires. This three-layer structure is modular: changing a component’s noise model is localised to a single object and does not require protocol logic to change.

### 3.1.2 Hardware Model

Physical impairments are introduced via a ten-layer parametric model. Each layer adds one impairment cumulatively to the previous, from a fully ideal baseline (layer 0) to a system incorporating all realistic imperfections (layer 9). This structure allows the marginal impact of each physical parameter on key rate to be isolated, forming the basis of the parameter sweep analysis in Section 4. Table 2 lists the ten layers.

| Layer | Parameter added          | Symbol          | Default value                        | Verification source           |
|-------|--------------------------|-----------------|--------------------------------------|-------------------------------|
| 0     | Ideal baseline           | —               | All perfect                          | —                             |
| 1     | Fibre attenuation        | $\alpha$        | 0.18 dB/km                           | Corning [31]                  |
| 2     | Detector efficiency      | $\eta_Z$        | 0.90                                 | ID Quantique [32]             |
| 3     | Dark count rate          | $d_c$           | 50 cps                               | ID Quantique [32]             |
| 4     | TX insertion loss        | $L_i$           | 0.10 (10%)                           | OZ Optics [33]                |
| 5     | RX node loss             | $\Delta_n$      | 2.0 dB                               | Duplinskiy <i>et al.</i> [34] |
| 6     | Source error rate        | $\varepsilon_s$ | 0.015                                | Quandela [35]                 |
| 7     | Fibre dephasing          | $\beta$         | $3.2 \times 10^{-7} \text{ km}^{-1}$ | Corning [31]                  |
| 8     | X-basis detector bias    | $\eta_X$        | 0.715                                | Grasselli <i>et al.</i> [36]  |
| 9     | Beam-splitter efficiency | $\eta_{BS}$     | 0.97                                 | Thorlabs [37]                 |
| —     | Relay position           | $\delta$        | 0.5                                  | User symmetry                 |

Table 2: Ten-layer hardware model. Layer  $k$  includes all impairments of layer  $k - 1$  plus the new parameter listed. Layer 9 is the realistic operating point used for network simulations. Empty cells in the verification column indicate parameters pending datasheet confirmation.

*Layer 0* is the ideal baseline: photons propagate without loss or noise, so key rate is limited only by propagation delay. Bipartite sweeps at this layer vary distance across the metropolitan range (0–50 km), establishing an upper bound on achievable key rate and confirming correct simulation timing before physical impairments are introduced.

*Layer 1* adds fibre attenuation: each photon is discarded before transmission with probability  $p_{\text{loss}} = 1 - 10^{-\alpha L/10}$ , where  $\alpha = 0.18 \text{ dB/km}$  and  $L$  is the arm length. For MDI-QKD the two arms attenuate independently, so coincidence probability scales as the product of both arm transmittances.

*Layer 2* introduces Z-basis detector efficiency  $\eta_Z = 0.90$ : each photon arriving at the receiver is accepted with probability  $\eta_Z$ . Missed detections reduce key length but do not introduce errors.

*Layer 3* adds dark count rate  $d_c = 50 \text{ cps}$ : a spurious click fires per time slot with probability  $d_c/(d_c + f)$ , where  $f$  is the source repetition rate. Dark clicks on empty slots generate false detections; dark clicks on occupied slots override the measurement with a random bit, contributing directly to QBER.

The default detector parameters ( $\eta_Z = 0.90$ ,  $d_c = 50 \text{ cps}$ , layers 2–3) correspond to a superconducting nanowire single-photon detector (SNSPD) operating point, verified against the ID Quantique (IDQ) ID281 [32]. The bipartite parameter sweeps in Section 4.1 additionally span the single-photon avalanche diode (SPAD) regime ( $\eta_Z \approx 0.10\text{--}0.25$ ,  $d_c \approx 100\text{--}200 \text{ cps}$ ), representative of the IDQ ID230 [38], covering the full range of detector technologies relevant to near-term metropolitan deployments.

*Layer 4* adds TX insertion loss  $L_i = 0.10$ : a fixed fraction of photons is lost at the source before entering the fibre, absorbed into the combined transmittance  $T = 10^{-\alpha L/10}(1 - L_i)$ . This models a polarisation-maintaining fused fibre coupler combining two orthogonally-polarised branches (e.g. H/V or  $\pm 45^\circ$ ) into a single output fibre for state preparation. Since the combined signals differ in polarisation, the coupler avoids the reciprocal 3 dB splitting loss that would apply to identical inputs [33]; only the excess loss is incurred. Taking the OZ Optics DTS0092 PM fused coupler’s maximum excess loss of 0.4 dB,  $L_i = 1 - 10^{-0.4/10} \approx 0.09$ , rounded up to 0.10 for a safety margin.

*Layer 5* introduces receiver node loss  $\Delta_n = 2.0$  dB as a single lumped parameter applied uniformly to both protocols, capturing the insertion loss of optical receiver components (EOM, PBS, HOM coupler). Each photon surviving fibre transmission is discarded with probability  $1 - 10^{-\Delta_n/10}$  before reaching the detector. This value matches the measured total insertion loss of Bob’s receiver stage (LiNbO<sub>3</sub> phase modulator for basis selection, followed by a polarising beamsplitter) in a demonstrated all-fibre polarisation-encoded BB84 system [34], giving  $\eta_{\text{node}} = 10^{-2.0/10} \approx 0.63$ . Individual component losses are not modelled separately and the protocol-specific breakdown is noted as a limitation in Section 5.

*Layer 6* adds source error rate  $\varepsilon_s = 0.015$ : each emitted photon receives an unintended  $\hat{X}$  flip with probability  $\varepsilon_s$ , modelling multi-photon emission events characterised by  $g^{(2)}(0)/2$  for the reference single-photon source [35]. The Quandela Prometheus datasheet specifies a second-order correlation  $g^{(2)}(0) < 0.03$ : the probability that a given pulse contains a multi-photon emission rather than a single photon. Taking the standard approximation that half of such multi-photon events go undetected as errors,  $\varepsilon_s \approx g^{(2)}(0)/2 < 0.015$ . Ideal single-photon sources are assumed throughout; WCP sources and decoy-state techniques [27] are not modelled, so reported rates represent an upper bound relative to practical laser deployments.

*Layer 7* adds fibre dephasing  $\beta = 3.2 \times 10^{-7} \text{ km}^{-1}$ : the `DephaseNoiseModel` applies a per-photon dephasing error with probability  $p_\phi = \min(1, \beta L)$  proportional to arm length, modelling polarisation-mode dispersion in standard single-mode fibre. This value is derived from the Corning SMF-28 PMD coefficient  $D_{\text{PMD}} = 0.04 \text{ ps}/\sqrt{\text{km}}$  [31]: mean differential group delay  $\Delta\tau = D_{\text{PMD}}\sqrt{L}$  induces a polarisation-dephasing probability  $p_\phi \approx (\Delta\tau/T)^2/2$ , with  $T = 50 \text{ ps}$  the detector timing-discrimination window (IDQ ID281 jitter [32]). Since  $\Delta\tau^2 \propto L$ ,  $p_\phi$  is linear in  $L$ , giving a constant per-km rate  $\beta = D_{\text{PMD}}^2/(2T^2)$ . Deployed fibre is assumed to vary  $\pm 20\%$  around this value ( $2.6 \times 10^{-7}$ – $3.8 \times 10^{-7} \text{ km}^{-1}$ ) to reflect manufacturing spread within the SMF-28 specification.

*Layer 8* introduces X-basis detector bias  $\eta_X = 0.715$ : photons measured in the X basis are detected with lower efficiency than those in the Z basis, reflecting the extra optical

stage (dispersive medium/interferometer) required for the test-basis measurement. This value follows Grasselli *et al.* [36], who model the test-basis (X) measurement's extra optical stage as an additional 1 dB insertion loss, giving  $\eta_X = \eta_Z \times 10^{-1/10}$ . Since insertion loss is a passive multiplicative factor applied upstream of the detector, this ratio is independent of the detector's own quantum efficiency and applies directly at  $\eta_Z = 0.90$ :  $\eta_X = 0.90 \times 10^{-1/10} \approx 0.715$ .

*Layer 9* (MDI only) adds beam-splitter efficiency  $\eta_{BS} = 0.97$ : the HOM coupler at Charlie's relay has non-unit splitting ratio, reducing coincidence detection probability per arm. This value is derived from the Thorlabs TN1550R5F2 fibre-coupled 50:50 coupler [37], a 1550 nm narrowband fused coupler specified with excess loss  $< 0.15$  dB:  $\eta_{BS} = 10^{-0.15/10} \approx 0.97$ . The relay position  $\delta = 0.5$  (Charlie at the link midpoint) is held fixed for all bipartite sweeps, reflecting user symmetry; it is determined geometrically in the network setting via Equation 8.

### 3.1.3 Bipartite Single-Link Analysis

The bipartite simulation sweeps one hardware parameter at a time, holding all others at their layer-9 defaults (Table 2). For each parameter value,  $r = 20$  independent Monte Carlo trials are run; each trial constitutes one complete QKD exchange from photon emission through key sifting. The mean key rate, mean QBER, and success rate (fraction of trials with QBER  $< 11\%$ ) are reported as a function of the swept parameter.

**BB84 simulation.** Alice and Bob are represented as `Node` objects connected by one quantum channel (Alice  $\rightarrow$  Bob) and two classical channels for basis announcement and sifting. Alice prepares  $n = 1,024$  photons at  $f = 10^7$  Hz. Each photon is assigned an independent random basis (Z or X) and bit value. Encoding is applied qubit-by-qubit: a bit-1 value applies  $\hat{X}$ ; an X-basis choice applies a Hadamard gate. Source imperfection ( $\varepsilon_s$ ) is modelled by an additional  $\hat{X}$  flip applied with probability  $\varepsilon_s$ , consistent with the multi-photon emission rate  $g^{(2)}(0)/2$  of the reference single-photon source. Photons are discarded before transmission with probability

$$p_{\text{loss}} = 1 - 10^{-\alpha L/10} (1 - L_i), \quad (1)$$

modelling both fibre attenuation and TX coupling loss. The quantum channel applies an additional per-photon dephasing error of probability  $\beta L$ . At Bob, each arriving photon is further subject to receiver node loss (probability  $1 - 10^{-\Delta_n/10}$ ), then detected with probability  $\eta_Z$  (Z basis) or  $\eta_X$  (X basis). Dark counts fire with per-slot probability  $d_c/(d_c + f)$ ; a dark click on a slot where a real detection also occurs overrides with a random

outcome. Dark clicks also fire on slots where the photon was lost.

After transmission, Bob announces his measurement bases and the indices of his detected slots; Alice returns her basis list. Bob retains only detections where both parties chose the same basis. The sifted key is the list of Bob's raw measurement outcomes at matching-basis slots. QBER is the fraction of bit disagreements between Alice's and Bob's sifted keys:

$$\text{QBER} = \frac{1}{\ell} \sum_{k=1}^{\ell} \mathbf{1}[a_k \neq b_k], \quad (2)$$

where  $\ell$  is the sifted key length. The QBER is measured over the full sifted key, not a sacrificed subset, since no eavesdropper is modelled. Key rate is then:

$$R = \frac{\ell}{T_{\text{end}} - T_{\text{start}}} \quad (\text{bps}), \quad (3)$$

where  $T_{\text{end}}$  is the simulation time at which Bob completes basis reconciliation and  $T_{\text{start}}$  is the photon emission start, both in nanoseconds. Trials with  $\text{QBER} > 11\%$  are excluded and counted as failures. No error correction or privacy amplification is applied;  $R$  therefore overestimates the secure key rate (see Section 5).

**MDI simulation.** Three nodes are instantiated: Alice, Bob, and Charlie. Two quantum channels (Alice  $\rightarrow$  Charlie, Bob  $\rightarrow$  Charlie) and six classical channels (BSM announcement and basis disclosure) are constructed. Each end node prepares  $n = 512$  photons — half the BB84 count — normalising comparison to a single-source rate scale, since MDI involves two independent emitters per trial. Encoding, source error, and per-photon transmission loss are applied identically to the BB84 case, with arm lengths  $L_A = L\delta$  and  $L_B = L(1 - \delta)$  where  $\delta = \text{charlie\_pos}$  locates Charlie on the link; each arm applies an independent dephasing model proportional to its length.

At Charlie, photons arriving from Alice and Bob on the same pulse index are paired. Both photons undergo node coupling loss ( $\Delta_n$ ) and are detected with probability  $\eta_{\text{BS}} \eta_d$  (HOM beam-splitter efficiency  $\times$  detector efficiency). A CNOT gate followed by a Hadamard on the control qubit and measurement of both qubits implements the Bell state measurement. The four computational outcomes correspond to the four Bell states:

$$(a, b) \longmapsto \begin{cases} |\Phi^+\rangle & a = 0, b = 0 \\ |\Phi^-\rangle & a = 1, b = 0 \\ |\Psi^+\rangle & a = 0, b = 1 \\ |\Psi^-\rangle & a = 1, b = 1 \end{cases} \quad (4)$$



Consistent with the linear-optical BSM constraint [10], only  $|\Psi^\pm\rangle$  outcomes ( $b = 1$ ) are treated as conclusive;  $|\Phi^\pm\rangle$  outcomes ( $b = 0$ ) are discarded. This yields a maximum BSM success probability of 50% for coincident pairs. Dark count clicks on unpaired or lost slots generate spurious conclusive outcomes, contributing to QBER. Charlie announces BSM outcomes and basis-mismatch indices to Alice and Bob; end users discard inconclusive and basis-mismatched rounds. Bob applies basis-dependent bit corrections according to the BSM outcome to restore key agreement with Alice. The sifted key, QBER, and key rate are computed as per Equations 2–3, with  $T_{\text{end}} = \max(T_{\text{end}}^{\text{A}}, T_{\text{end}}^{\text{B}})$ . In the network setting, the Charlie position is determined geometrically by the relay’s location relative to the user pair (Section 3.1.4).

### 3.1.4 Multi-User Network Analysis

The multi-user simulation models a metropolitan network of  $N$  users within an  $A \times A$  km service area. Users are placed uniformly at random within the grid.

**Relay placement.** Placing the  $K$  MDI relay nodes is a facility-location problem: given fixed user locations and free relay nodes that connect to their assigned users and, for  $K > 1$ , to each other over a fully meshed backbone (Table 3), the placement that minimises deployment cost is the one that minimises total fibre length, since fibre enters the cost model linearly through  $c_f$  (Equation 10, Section 3.2). Two closed-form single-facility results bound the problem. Minimising the sum of *squared* distances from a set of fixed points to one free node is solved by the centroid (the mean position), the standard least-squares result obtained by zeroing the gradient of a sum of quadratics. Minimising the sum of *unsquared* Euclidean distances instead is solved by the geometric median (the Weber point), which has no closed form in general and is instead obtained iteratively via Weiszfeld’s algorithm,

$$x^{(t+1)} = \frac{\sum_i p_i / \|x^{(t)} - p_i\|}{\sum_i 1 / \|x^{(t)} - p_i\|}, \quad (5)$$

iterated to convergence. The two solutions coincide only for symmetric configurations. Since total fibre length is a sum of *unsquared* distances, the geometric median — not the centroid — is the objective that actually matches the network’s linear cost structure; the centroid instead minimises mean squared path length, a quantity relevant to statistical clustering (and to path loss under a squared-distance model) but not to CapEx.

For  $K > 1$  fully-meshed relays, the placement problem does not decompose into  $K$  independent per-cluster geometric medians, because the inter-relay backbone term couples all relay positions together. Writing cluster  $k$ ’s user set as  $C_k$  and its relay position as  $x_k$ ,

total fibre length is

$$L(x_1, \dots, x_K) = \underbrace{\sum_{k=1}^K \sum_{i \in C_k} \|x_k - p_i\|}_{\text{spoke cost}} + \underbrace{\sum_{k < l} \|x_k - x_l\|}_{\text{backbone (mesh) cost}}. \quad (6)$$

This is the classical multi-facility Weber problem, dating to Miehle’s original network-of-junctions formulation [39] and specialised to hub-and-spoke interconnection by O’Kelly [40]; Alumur and Kara [41] survey the broader hub-location literature that has since grown around it.  $L$  is jointly convex in  $(x_1, \dots, x_K)$ , being a sum of norms, so a unique global minimum exists; but setting each  $x_k$  independently to its own cluster’s geometric median ignores the gradient contribution of the mesh term and is generally suboptimal. Zeroing  $\nabla_{x_k} L$  gives the generalised (multi-facility) Weiszfeld fixed-point update, with convergence to the joint optimum guaranteed by convexity [42]:

$$x_k^{(t+1)} = \frac{\sum_{i \in C_k} \frac{p_i}{\|x_k^{(t)} - p_i\|} + \sum_{l \neq k} \frac{x_l^{(t)}}{\|x_k^{(t)} - x_l^{(t)}\|}}{\sum_{i \in C_k} \frac{1}{\|x_k^{(t)} - p_i\|} + \sum_{l \neq k} \frac{1}{\|x_k^{(t)} - x_l^{(t)}\|}}, \quad (7)$$

updating all  $K$  relays simultaneously each iteration.

This work places relays by the generalised Weiszfeld iteration of Equation 7, in place of a plain  $k$ -means centroid. Cluster assignment (spoke membership  $C_k$ ) is still initialised by  $k$ -means with  $n_{\text{init}} = 10$  random restarts, a reasonable heuristic for partitioning users spatially; relay *position* within each cluster is then refined by Equation 7 rather than fixed at the cluster centroid, and assignment is re-evaluated against nearest-relay distance after each position update, alternating position refinement and reassignment to convergence. For  $K = 1$  this reduces to the plain single-facility Weiszfeld iteration of Equation 5, since the mesh term is absent; the joint coupling in Equation 7 has an effect only for  $K > 1$ . User-to-relay assignment remains `user_relay[i] = arg min_k ||u_i - r_k||_2`, evaluated against the converged relay positions. This is a methodological change from earlier  $k$ -means-centroid placement used elsewhere in this work’s development. It is used exactly as described here —  $k$ -means-seeded, Weiszfeld-refined, with membership re-evaluated by nearest-relay distance after every position update — specifically for the relay count sweep of Section 3.3, where relay infrastructure is genuinely redesigned from scratch at each  $K$  (a greenfield deployment decision). The user count sweep of Section 3.3 and the relay-placement strategy comparison of Appendix B instead use the same position

update (Equation 7) with catchment membership held *fixed* rather than reassigned: those experiments specifically test how well a placement decision made once serves users it was never re-optimised for, which the reassigning variant would undermine by silently re-clustering around whatever the relay’s current position happens to be. Appendix B.4 justifies that choice in detail. The geometric-median objective could still revert to a plain centroid in either variant if the fibre-length cost model is later judged not to warrant the added iteration.

To decouple topology variability from Monte Carlo noise, all network experiments are repeated across  $S$  independently drawn user placements (seeds), with relay positions re-optimised per seed. Aggregate statistics across seeds produce error bars that characterise topology uncertainty — the variance arising from user placement — rather than the sampling noise of a fixed topology. The number of seeds is chosen such that the inter-seed standard deviation of the protocol comparison metric (mean key rate) is small relative to the protocol gap.

This work focuses on metropolitan-scale networks, where inter-node distances are short enough for both BB84 and MDI-QKD to establish keys at viable rates. Two network protocols are simulated:

**BB84 mesh.** Every user pair  $(i, j)$  communicates over a dedicated direct dark-fibre link. Per-link tortuosity is sampled independently for each of the  $N(N - 1)/2$  links:  $\tau_{ij} \sim \max(1, \mathcal{N}(\mu_\tau, \sigma_\tau^2))$ . The physical cable length is  $d_{ij} \tau_{ij}$  km. All pairs are dispatched to the multiprocessing pool simultaneously;  $r$  Monte Carlo trials are run per pair. Network key rate is the mean over all valid (pair, trial) combinations.

**MDI network.** Tortuosity is sampled once per physical cable rather than per pair, ensuring that all simulation pairs sharing a cable see the same physical routing factor. Specifically,  $N$  user-to-relay cables and  $K(K - 1)/2$  relay-to-relay backbone cables each draw one independent tortuosity factor at the start of each seed.

For a same-cluster pair  $(i, j)$  assigned to relay  $R_i = R_j = R$ , Alice’s arm length is  $d_{iR} \tau_{iR}$  and Bob’s arm length is  $d_{jR} \tau_{jR}$ ; the BSM takes place at relay  $R$ . For a cross-cluster pair  $(i, j)$  with  $R_i \neq R_j$ , Bob’s photon travels from user  $j$  to relay  $R_j$ , then along the backbone to relay  $R_i$  where Alice’s BSM station is located. Bob’s effective arm length is  $d_{jR_j} \tau_{jR_j} + d_{R_j R_i} \tau_{R_j R_i}$ . The Charlie position parameter is set geometrically:

$$\delta = \frac{d_{iR_i} \tau_{iR_i}}{d_{iR_i} \tau_{iR_i} + d_{jR_j} \tau_{jR_j} + d_{R_j R_i} \tau_{R_j R_i}}. \quad (8)$$

An optical switch at each relay routes backbone photons to the BSM station for cross-cluster pairs. Its insertion loss is set to 0 dB in the current model; the asymmetry between

switch installation cost (per relay) and switch-induced loss (cross-cluster pairs only) is noted as a simulation limitation in Section 5.

The network key rate reported throughout this work is the mean pairwise point-to-point key rate averaged over all  $N(N - 1)/2$  user pairs and  $S$  topology seeds. This is not a live traffic simulation: no channel contention, scheduling, or multiplexing is modelled. Every user pair is assumed to extract key simultaneously and independently at full channel capacity, with no interference from other pairs sharing physical infrastructure. The reported metric therefore constitutes a semi-probabilistic geometric upper bound on network throughput: *geometric* because the rate for each pair is determined by the physical path geometry (distances, relay placement, tortuosity); *probabilistic* because relay positions and cable tortuosities are drawn stochastically and averaged across seeds; and an *upper bound* because real deployments would see contention for shared fibre, detector dead time across concurrent channels, and scheduling overhead that this model does not capture. The multi-seeding procedure described above quantifies the geometric component of this bound across the distribution of plausible topologies for a given  $N$ ,  $K$ , and service area.

## 3.2 Cost Model

We model the total deployment cost of a QKD network algebraically using symbolic per-component parameters. All costs are capital expenditure (CapEx) only; operational expenditure — in particular the cryogenic cooling required by superconducting nanowire single-photon detectors (SNSPDs) and quantum dot (QD) sources — is not modelled. Classical communication infrastructure (basis reconciliation, error correction, and key forwarding) is assumed to be a sunk cost of existing network infrastructure. The most directly related prior cost analysis is that of Karavias *et al.* [22], which applies integer linear programming to PM-QKD and entanglement-based infrastructure planning with concrete component costs; MDI-QKD is not considered there. We extend this structural analysis to MDI-QKD, deriving cost formulae that expose the scaling difference between the two protocols without requiring absolute price data.

### 3.2.1 Component Counts

The two network topologies impose fundamentally different hardware structures. In BB84, every user pair requires a dedicated direct dark-fibre link: the network is a complete mesh of  $N(N - 1)/2$  quantum channels. At each receiver, a polarising beam splitter (PBS) separates Z and X basis outputs to two single-photon detectors (SPDs), and an electro-

optic modulator (EOM) performs active basis selection. No relay infrastructure is required.

In MDI-QKD, users transmit photons to BSM relay stations rather than to each other. Each relay is a passive station comprising a 50:50 HOM beam splitter, two PBS (one per arm), four SPDs, and one optical switch for routing photons from cross-cluster pairs. The network requires  $N$  user-to-relay fibre links and  $K(K-1)/2$  relay-to-relay backbone links.

Table 3 summarises the component counts for both protocols.

| Component                 | BB84               | MDI                    |
|---------------------------|--------------------|------------------------|
| Single-photon sources     | $N$                | $N$                    |
| Single-photon detectors   | $2N$               | $4K$                   |
| 50:50 beam splitter (HOM) | 0                  | $K$                    |
| Polarising beam splitter  | $N$                | $2K$                   |
| Electro-optic modulator   | $N$                | 0                      |
| Optical switch            | 0                  | $K$                    |
| Fibre links               | $\frac{N(N-1)}{2}$ | $N + \frac{K(K-1)}{2}$ |

Table 3: Hardware component counts for a network of  $N$  users and  $K$  relay nodes. MDI relay nodes are passive BSM stations. BB84 forms a direct mesh with no relay infrastructure.

Fibre cost scales as  $\mathcal{O}(N^2)$  for BB84 and as  $\mathcal{O}(N)$  for MDI at fixed  $K$ , providing the structural cost incentive for relay-based architectures as user counts grow.

### 3.2.2 Symbolic Cost Formulae

Manufacturer prices for quantum dot photon sources and single-photon detectors are not publicly available. We therefore parameterise the cost model with three symbolic constants:  $c_s$  (cost per photon source),  $c_d$  (cost per single-photon detector), and  $c_f$  (cost per km of installed dark fibre). Using the component counts in Table 3, the total CapEx for each protocol is:

$$C_{\text{BB84}} = N c_s + 2N c_d + \frac{N(N-1)}{2} \bar{d}_{\text{BB84}} c_f, \quad (9)$$

$$C_{\text{MDI}} = N c_s + 4K c_d + \left( N \bar{d}_u + \frac{K(K-1)}{2} \bar{d}_r \right) c_f, \quad (10)$$

where  $\bar{d}_{\text{BB84}}$  is the mean user-pair distance in the BB84 mesh,  $\bar{d}_u$  the mean user-to-relay distance, and  $\bar{d}_r$  the mean relay-to-relay backbone link length. Passive relay hardware (HOM beam splitters, polarising beam splitters, optical switches) contributes additional fixed cost per relay and is omitted here for clarity.

The dominant asymptotic difference between the two protocols is the fibre term: BB84 scales as  $\mathcal{O}(N^2)$  (one dedicated link per user pair) whilst MDI scales as  $\mathcal{O}(N)$  at fixed relay count  $K$ . Setting  $C_{\text{BB84}} = C_{\text{MDI}}$  and solving for  $N$  yields a crossover threshold  $N^*$  above which MDI is cheaper, whose value depends on the ratio  $c_d/c_f$ , relay count  $K$ , and network geometry.

Commercial SNSPD detectors are typically sold as multi-channel units integrating several channels within a single cryostat (e.g. 2–8 channels per unit), so an MDI relay’s four SPDs may in practice be equipped with a single packaged unit rather than four independent  $c_d$  purchases; this packaging effect is not reflected in Equation 10, and its implications for the cost crossover are discussed in Section 4.3.

In practice, installed cable follows duct routes longer than the straight-line Euclidean distance. We model this through per-link tortuosity: each physical cable  $(i, j)$  is assigned an independent factor  $\tau_{ij}$  drawn from a truncated normal distribution:

$$\tau_{ij} \sim \max(1, \mathcal{N}(\mu_\tau, \sigma_\tau^2)), \quad \mu_\tau = 1.2, \quad \sigma_\tau = 0.1. \quad (11)$$

The mean  $\mu_\tau = 1.2$  is representative of urban cable routing, where installed fibre is typically 15–25% longer than the Euclidean distance. Tortuosity affects the simulated key rate through increased fibre loss on longer cables. Each physical cable draws one factor, held constant across all simulation runs; for MDI, user-to-relay and relay-to-relay cables each draw a single factor shared across all user pairs that traverse that cable, ensuring physical consistency when the same cable appears in multiple paths.

### 3.3 Scalability Analysis

Scalability is characterised through two orthogonal sweeps.

**User count sweep.** Rather than generating the full  $N_{\text{max}}$ -user network up front and revealing a prefix of it, users are added one at a time from a single random stream: each new user joins a uniformly-random existing relay catchment and is placed by an isotropic Gaussian around that catchment’s centre, so that  $N \rightarrow N + 1$  genuinely means *adding* a user rather than redrawing an unrelated sample. Relay positions are placed once, for the  $N_{\text{min}}$ -user network, by whichever placement strategy is specified (centroid, boundary, or the Weiszfeld iteration of Equation 7), and then held fixed as  $N$  grows to  $N_{\text{max}}$ ; catchment membership, once assigned, is likewise never reassigned by nearest-relay distance. This models a provider that sizes relay infrastructure for an initial deployment and serves subsequent organic growth without continuously re-optimising it, rather than a provider that instantly re-plans infrastructure after every new customer. For each  $N$ ,

both protocols are simulated and their mean key rate is recorded. The marginal rate  $\Delta\bar{R}/\Delta N$  is computed numerically from adjacent sweep points. The algebraic cost model separately predicts that MDI marginal CapEx decreases with  $N$  (relay hardware is sunk, only user-to-relay fibre grows linearly), whilst BB84 marginal CapEx grows as  $\mathcal{O}(N)$  since each new user requires  $N - 1$  new direct fibre links.

**Relay count sweep.** User count  $N$  is fixed and  $K$  is varied from 1 to  $K_{\max}$ . BB84 is relay-independent and appears as a constant baseline. MDI key rates improve as  $K$  increases (shorter mean user-to-relay distances) until the relay-to-relay backbone becomes the dominant bottleneck. This sweep surfaces the performance-optimal relay count for a given  $N$  and area.

## 4 Results

### 4.1 Point-to-Point Performance

Eleven independent parameter sweeps characterise bipartite BB84 and MDI-QKD key rate around a common metropolitan operating point: 20 km Alice–Bob separation, 0.18 dB/km fibre attenuation,  $\eta_d = 0.9$  detector efficiency, 50 cps dark count rate,  $L_i = 0.1$  TX insertion loss, 2.0 dB node loss,  $\varepsilon_s = 0.015$  source error rate,  $\beta = 3.2 \times 10^{-7}$ /km fibre dephasing, and  $\eta_{bs} = 0.97$  beam-splitter efficiency, with a single parameter varied at a time and all others held at this operating point. Each point is the mean secure key rate over 1000 Monte Carlo runs; shaded bands show the sample spread. Four sweeps that isolate the dominant geometric and hardware effects — distance, detector efficiency, receiver (node/connector) loss, and Charlie’s relay position — are discussed here; the full set of eleven, including these four, is reproduced in Appendix A.

**Distance.** Figure 4 sweeps Alice–Bob separation from 1–100 km at fixed attenuation.




Figure 4: Absolute and relative secure key rate versus Alice–Bob separation, at fixed 0.18 dB/km fibre attenuation. The PLOB bound is overlaid on the left axis. Shaded vertical bands mark the metropolitan (5–50 km) and long-range (> 80 km) regimes.



The BB84/MDI ratio stays close to a factor of five throughout:  $4.3\text{--}5.5\times$  over the full 1–100 km range, varying by no more than  $\pm 15\%$  even as the underlying key rate itself falls by close to three orders of magnitude — Figure 7 below explains why this ratio is so stable. Both protocols also fall well below the PLOB bound, and increasingly so with distance: at 1 km, BB84 and MDI sit a factor of  $5.2\times$  and  $22.5\times$  below PLOB respectively; by 100 km this has widened to  $43\times$  and  $204\times$ . This growing shortfall is the practical motivation for the range-extending architectures discussed in Section 5.1 (*Metropolitan network scope*) and revisited as future work in Section 6.

**Detector- and receiver-side parameters.** Figure 5 sweeps detector efficiency  $\eta_d$  from 1.0 down to 0.15, shaded to mark the efficiency bands typical of SNSPD (0.80–0.95, IDQ ID281) and SPAD (0.10–0.25, IDQ ID230) technology, and Figure 6 sweeps receiver node/connector loss from 0–6 dB.




Figure 5: Absolute and relative secure key rate versus detector efficiency  $\eta_d$ , at fixed 20 km separation and 0.18 dB/km attenuation. Shaded bands mark typical SNSPD and InGaAs SPAD operating ranges.




Figure 6: Absolute and relative secure key rate versus receiver node/connector loss, at fixed 20 km separation and 0.18 dB/km attenuation. The PLOB bound is overlaid on the left axis; the shaded band marks the realistic hardware regime (0.5–4.0 dB).

Unlike distance, the BB84/MDI ratio is not stable under either of these sweeps: under detector efficiency it grows from  $4.7\times$  at  $\eta_d = 1.0$  to  $25.4\times$  at  $\eta_d = 0.15$ , a 5.4-fold widening across the sweep; under receiver loss it grows from  $3.3\times$  at 0 dB to  $13.1\times$  at 6 dB, a 4.0-fold widening. This is the direct empirical signature of parameters that act *per arm*, after Alice’s and Bob’s photons have already diverged towards Charlie: detector efficiency and receiver optics loss both enter MDI’s coincidence probability once per arm (quadratically in total) but BB84’s single-photon detection probability only once (Section 5.1), so a degradation in either parameter degrades MDI disproportionately — the same mechanism, acting on two different pieces of hardware. This has direct practical implications. The choice between SNSPD and InGaAs SPAD detectors matters far more for an MDI deployment than for a BB84 deployment of otherwise identical geometry. It also sharpens the caveat raised in Section 5.3 (*Receiver node loss*): since the model applies a single lumped  $\Delta_n = 2.0$  dB uniformly to BB84’s Bob-side optics and MDI’s Charlie-side optics, and MDI’s true receiver components (HOM beam splitter, PBS, optical switch) plausibly carry higher insertion loss than BB84’s simpler EOM/PBS chain, this sweep’s

widening ratio should be read as a lower bound on how much receiver hardware choice actually separates the two protocols.

**Relay placement.** Figure 7 sweeps Charlie’s position as a fraction of a fixed 20 km Alice–Bob distance.




Figure 7: MDI key rate versus Charlie’s position as a fraction of the fixed 20 km Alice–Bob distance, with BB84 plotted as a distance-independent reference. The shaded band marks the near-symmetric regime.

MDI key rate peaks at the symmetric midpoint (61.7 kbps at 0.5) and falls by over a third towards either end (39.0 kbps at 0.1; 39.2 kbps at 0.9), confirming that Charlie sitting at the arm midpoint is rate-optimal. This is what makes the distance ratio above so stable: with each arm carrying half of the total channel loss at that optimum, the two-photon coincidence probability  $\eta_{\text{arm}}^2 \eta_{bs} \eta_d^2$  scales with total link transmittance in step with BB84’s single-photon term  $\eta_{\text{link}} \eta_d$  (Section 5.1), so channel-side loss scales both protocols together rather than opening up the gap between them. Moving Charlie off-centre breaks this cancellation and the BB84/MDI ratio widens to 8.2–8.3 $\times$  at the sweep extremes, close to a 60% increase over the midpoint ratio of 5.2 $\times$  — a clean illustration that relay geometry, and not only raw hardware quality, materially affects the security–rate trade-off.

Together, these four sweeps indicate that within the metropolitan regime studied here, the BB84/MDI key rate gap is set primarily by a roughly constant structural factor of  $\sim 5\times$  arising from the two-photon BSM requirement and Charlie’s optimal midpoint placement, and that per-arm hardware parameters — detector efficiency and receiver loss, not link distance or relay geometry — are what is most likely to widen this gap substantially in practice. Three further views of the same operating point — the cumulative hardware layers of Table 2 swept against distance, a fixed-distance ranking of all nine parameters by relative impact, and the isolated per-parameter distance sweeps that connect the two — test this conclusion directly.

**Cumulative hardware layers.** Figure 8 sweeps distance for each of the ten cumulative hardware layers of Table 2 in turn, rather than isolating one parameter at a time as in the sweeps above.




Figure 8: Secure key rate versus Alice–Bob separation under each cumulative hardware layer of Table 2, from the ideal baseline (layer 0, darkest) to the full realistic model (layer 8 for BB84, layer 9 for MDI, lightest). Each layer includes all impairments of the previous one plus the parameter named in its legend entry. Error bars show the Monte Carlo sample spread.

Both protocols show the same qualitative shape: a large step from layer 0 to layer 1 followed by a sequence of comparatively small further steps. For BB84 the layer 0-to-1 step (adding fibre attenuation alone) is barely visible at 1 km — 4447.7 kbps falls to 4274.8 kbps, a 3.9% reduction — but dominates entirely by 100 km, where the same step takes the rate from 374.4 kbps to 5.9 kbps, a 98.4% reduction. MDI shows the identical pattern: 3.4% at 1 km widening to 98.5% at 100 km. Fibre attenuation is therefore not a fixed per-parameter penalty in the way the other eight layers are; its fractional impact is itself a

function of distance, negligible at short range and overwhelming at long range, consistent with it being the sole layer whose loss probability scales with  $L$  directly (Section 3.1.2) rather than acting as a fixed multiplicative factor at any given operating point. Against this backdrop, layer 5 (node loss) stands out among the remaining eight as the one further step large enough to see clearly even at short range: for BB84 the layer 4-to-5 step at 1 km takes 3653.3 kbps to 2310.2 kbps, a 36.8% reduction; for MDI, 1508.8 kbps to 598.1 kbps, a 60.4% reduction. Both figures are consistent with the sensitivity ranking below, and both closely match the isolated node-loss impact reported next, despite being measured from an entirely different sweep (cumulative layer stacking rather than single-parameter isolation).

**Parameter sensitivity ranking.** Figure 9 ranks all nine layer-9 hardware parameters directly by their individual contribution to key-rate loss, at a fixed 25 km operating point (ideal key rate 1223.85 kbps for BB84, 482.15 kbps for MDI).




Figure 9: Percentage key-rate reduction from the ideal (layer-0) baseline at 25 km separation, for each parameter set individually to its layer-9 realistic value (Table 2). BS efficiency applies to MDI only; the BB84 bar is fixed at zero.

Fibre attenuation dominates both protocols almost identically at this fixed distance (64.5% BB84, 64.3% MDI), the two curves not yet having diverged the way they do by 100 km in Figure 8. Every other parameter instead reproduces the per-arm asymmetry established in Section 4.1: node loss reduces BB84 by 36.7% but MDI by 60.5%; basis bias by 14.1% versus 24.5%; TX insertion loss by 10.0% versus 19.4%; detector efficiency by 4.8% versus 9.4% — each roughly 1.6–2 $\times$  more damaging to MDI than to BB84, the signature of a parameter acting once per arm rather than once per link. Beam-splitter

efficiency, applicable to MDI alone, contributes a further 6.0%. Dark count rate, source error rate, and fibre dephasing are negligible for both protocols at this operating point (under 1% in every case), matching the "noise-like parameters" grouping of the appendix sweeps (Appendix A).

**Isolated parameter impact.** Figure 10 extends the fixed-distance ranking above across the full 1–100 km range, setting exactly one parameter to its layer-9 value at a time against an otherwise fully ideal baseline.




Figure 10: Secure key rate versus Alice–Bob separation with exactly one hardware parameter set to its layer-9 realistic value (Table 2) and all others held ideal, against an ideal (baseline) reference. Error bars show the Monte Carlo sample spread.

The two parameters singled out above behave in structurally different ways across distance. Fibre loss (blue) tracks the ideal baseline closely at short range and then pulls away with distance — for BB84, a 4.3% reduction at 1 km widening to 98.4% at 100 km; for MDI, 3.7% widening to 98.2% — confirming, from an independent sweep, the same distance-dependence seen in Figure 8. Node loss instead reduces key rate by an essentially *constant fraction at every distance*: 36.9% at 1 km versus 36.8% at 100 km for BB84, and 60.3% at both distances for MDI, matching the 25 km sensitivity figures (36.7% and 60.5%) to within Monte Carlo noise. This distinction reflects the two parameters' underlying physics: fibre loss is proportional to path length, so its effect is negligible over short arms and dominant over long ones, whereas node loss is a fixed per-arm insertion loss applied once regardless of distance, so its fractional impact does not change as the link is extended. Node loss, basis bias, TX insertion loss, and (for MDI) beam-splitter efficiency therefore constitute a qualitatively distinct category of impairment from fibre

attenuation: fixed-fraction, per-arm penalties that erode key rate even over the shortest simulated separations and that consistently favour BB84 over MDI by the  $\sim 1.6\text{--}2\times$  margin identified throughout this section, rather than range-dependent losses that only become significant as the link is extended.

## 4.2 Network Performance

The user-count and relay-count sweeps of Section 3.3 are reported together here, since both characterise network-level key rate as a function of the same underlying variables ( $N$  and  $K$ ) rather than testing independent hypotheses.

## 4.3 Cost

The cost formulae of Section 3.2 predict that the fibre term dominates at large  $N$  for any positive  $c_f$ , so a crossover user count  $N^*$  above which MDI is cheaper always exists; the practical question is whether  $N^*$  falls within the user counts of realistic metropolitan deployments, which depends on the ratio  $c_d/c_f$ , relay count  $K$ , and network geometry.

Commercial detector packaging complicates this picture. SNSPDs are typically sold as multi-channel units integrating several channels within a single cryostat (e.g. 2–8 channels per unit), so an MDI relay’s four SPDs may in practice be equipped with a single packaged unit rather than four independent  $c_d$  purchases, reducing per-relay detector cost below  $4c_d$ . This narrows MDI’s cost advantage at small  $N$  and shifts  $N^*$  upward; since it cannot be quantified without product pricing, it is treated here as a sensitivity consideration on the crossover result rather than a correction to Equation 10.

## 4.4 Deployment Recommendations

Protocol selection recommendations follow from combining the simulation key rate results with the algebraic cost model. The relay sweep provides mean key rate  $\bar{R}(K)$  at fixed  $N$  and area; Equations 9–10 express cost as a function of the same variables and the free parameters  $c_s$ ,  $c_d$ ,  $c_f$ .

For a provider with a minimum key rate requirement  $R_{\min}$ , the relay sweep identifies the minimum relay count  $K^*$  satisfying  $\bar{R}(K^*) \geq R_{\min}$ . The corresponding MDI cost  $C_{\text{MDI}}(N, K^*, c_s, c_d, c_f)$  can then be compared to the BB84 cost  $C_{\text{BB84}}(N, c_s, c_d, c_f)$  symbolically, revealing the conditions on  $c_d/c_f$  under which MDI’s security advantage is obtained within the same budget as BB84. The crossover user count  $N^*$  derived from the cost formulae provides a scale threshold independent of absolute component prices.

## 5 Discussion

The results presented in Section 4 support a consistent picture of the trade-off space between BB84 and MDI-QKD. We discuss the key findings, their physical interpretation, the limitations of the methodology, and their implications for protocol selection.

### 5.1 The Security–Rate Trade-off

MDI-QKD’s elimination of detector side-channel attacks comes at a key rate penalty rooted in its physical architecture. Whereas BB84 requires one photon to traverse the full link and be detected, MDI requires two photons — one from each party — to arrive at Charlie’s relay simultaneously and produce a coincidence event. Loss on each arm compounds multiplicatively: if the single-arm transmittance is  $\eta_{\text{arm}}$ , the BSM detection probability scales as  $\eta_{\text{arm}}^2 \eta_{\text{BS}} \eta_d^2$ , compared to  $\eta_{\text{link}} \eta_d$  for BB84. At metropolitan distances and typical detector efficiencies, this effectively squares the channel transmittance, consistent with the order-of-magnitude rate gap between the two protocols in the bipartite setting. This is borne out empirically by the point-to-point sweeps in Section 4.1: the BB84/MDI ratio stays close to a constant  $\sim 5\times$  under channel-side sweeps (distance, attenuation) but widens to over  $25\times$  under a detector efficiency sweep (Figure 5), the signature of a parameter that acts quadratically on MDI’s coincidence probability but only linearly on BB84’s single-photon detection probability. The sensitivity ranking in Figure 9 quantifies this split directly at a fixed 25 km operating point: fibre attenuation degrades both protocols by almost the same fraction (64.5% BB84, 64.3% MDI), while every per-arm parameter — node loss, basis bias, TX insertion loss, detector efficiency, and (MDI-only) beam-splitter efficiency — degrades MDI by roughly  $1.6\text{--}2\times$  more than BB84 in relative terms. The cumulative-layer and isolated-parameter sweeps of Figure 8 and Figure 10 show why the two channel-side and per-arm categories behave so differently with distance: fibre attenuation’s fractional impact grows from a few percent at 1 km to over 98% at 100 km for both protocols, since loss probability scales with path length directly, whereas node loss removes an essentially constant fraction of key rate — around 37% for BB84 and 60% for MDI — at every distance from 1–100 km, since it is a fixed per-arm insertion loss applied once regardless of link length.

At the network level, the rate gap between MDI and BB84 widens relative to the bipartite case. MDI users route photons to a relay rather than directly to their communication partner, incurring a geometric detour for cross-cluster pairs. This routing overhead is absent in BB84’s direct mesh. The relay distance advantage is a function of relay count



$K$  and user distribution; the relay sweep quantifies the extent to which increasing  $K$  can recover the MDI rate penalty.

A third trust model is available in the simulation infrastructure: trusted-node BB84 (TBB84), implemented in `trusted_bb84_network.py` and described in Section 3.1.4. TBB84 routes a key between non-adjacent users via one or more relay hops, with each hop running standard BB84 and the resulting per-hop keys classically XOR-combined along the path to produce an end-to-end key. This trust model is fundamentally different from MDI’s. MDI’s Bell-state relay discloses only the joint correlation between Alice’s and Bob’s photons; even a fully adversarial Charlie learns nothing about either party’s individual bit values, so no trust in the relay is required. A TBB84 relay, by contrast, handles plaintext-equivalent key material at every hop it forwards: if a relay node is physically compromised, the attacker recovers the full end-to-end key for any pair routed through it. TBB84’s security therefore rests on the physical security of relay hardware, whereas MDI’s rests on detector integrity — the two protocols defend against different adversaries rather than offering directly comparable guarantees.

Because TBB84’s principal advantage over direct BB84 is range extension, obtained precisely by accepting this relay-trust assumption, it is only exercised in this work’s real-topology experiments (Section 3.1.4), where relay hops span distances that a single link could not otherwise reach. Within the  $10 \times 10$  km synthetic metropolitan grid used for the core BB84-vs-MDI comparison, direct BB84 links and MDI’s Bell-state relay both already achieve useful key rates without range extension (Section 5, *Metropolitan network scope*), so TBB84 would offer no rate or reach advantage there — only the strictly weaker trust assumption of trusted-node relaying, for no compensating benefit at this scale. TBB84 is therefore excluded from the main relay-count and user-count sweeps and reserved for deployments where relay range-extension is structurally required.

## 5.2 Cost Model Limitations

The cost model is a first-order CapEx-only algebraic estimate. Several factors are not captured.

*Component price data.* Manufacturer prices for quantum dot photon sources and SNSPD detectors are not publicly available; academic literature reports a wide range of estimates, and commercial quotes are required for quantitative comparison. The symbolic formulation (Equations 9–10) defers this uncertainty to  $c_s$ ,  $c_d$ , and  $c_f$ , but the sensitivity of the crossover  $N^*$  to these parameters is significant: a high  $c_d/c_f$  ratio (expensive detectors relative to fibre) widens the advantage of BB84’s lower detector count per user, whilst a high  $c_f/c_d$  ratio amplifies MDI’s  $\mathcal{O}(N)$  vs  $\mathcal{O}(N^2)$  fibre advantage.

*Operational expenditure.* Cryogenic systems — both SNSPD detectors and QD sources — require continuous cooling to temperatures of order 2–4 K, incurring energy costs that can rival hardware CapEx over a ten-year deployment lifetime [23]. The energetic cost of operating  $K$  SNSPD-equipped MDI relay stations may be significant at scale and is not captured by the CapEx-only model.

*BB84 fibre scaling at large  $N$ .* The BB84 direct mesh assumes one independent dark-fibre run per user pair. In practice, WDM or switched architectures would serve many pairs over shared fibre infrastructure at high  $N$ , substantially reducing BB84 fibre cost. The  $\mathcal{O}(N^2)$  fibre scaling should therefore be interpreted as an upper bound applicable to passive, unmanaged dark-fibre mesh deployments; it is physically exact for that specific case.

### 5.3 Simulation Model Limitations

*Source brightness and insertion loss decoupling.* The Quandela Prometheus single-photon source used as the reference product operates at 975 nm and achieves a minimum fibred brightness of 40%, with frequency conversion required to reach the 1550 nm telecom band. In the simulation, source brightness is absorbed into the effective source repetition rate (`source_freq` =  $10^7$  Hz, consistent with the 32 MHz raw rate after 40% brightness), and the `init_loss` parameter models separate optical path losses independently. Frequency conversion efficiency, additional conversion-stage noise, and the explicit coupling between source brightness and photon rate are not separately modelled. This is a deliberate simplification; a full model would propagate each loss stage independently and account for Raman background photons introduced by the conversion stage.

*Ideal single-photon sources.* All simulations assume one photon per pulse with no multi-photon component. WCP laser sources are overwhelmingly more practical and affordable than QD sources, and their multi-photon components require decoy-state analysis [27] to extract secure key rate. The ideal source assumption gives an upper bound on key rate and a lower bound on source cost; both should be re-evaluated for WCP deployments.

*Asymptotic key rates.* Finite-key corrections become significant at low photon count rates and short block lengths, reducing the extractable key compared to the asymptotic rate reported here. At long distances where photon detection events are sparse, the finite-key penalty is expected to be most severe for MDI-QKD, potentially widening the gap with BB84 beyond what the asymptotic analysis suggests.

*Relay placement.*  $k$ -means relay placement minimises mean user-to-relay distance but is not provably optimal for network key rate or cost. ILP-based placement [22] could provide tighter lower bounds on minimum MDI deployment cost, and would be a valuable

extension of this work.

*Receiver node loss.* The model applies a single lumped node loss  $\Delta_n = 2.0$  dB uniformly to both BB84 and MDI-QKD receivers. In reality, the optical components at each receiver are protocol-specific: BB84 uses an EOM and PBS, whilst MDI arms use a HOM beam splitter, PBS, and (for cross-cluster routing) an optical switch. These components carry different insertion losses and are purchased and maintained separately. A rigorous model would assign per-component losses derived from product datasheets, track them independently in the cost model, and apply the correct protocol-specific total at each receiver. This is left as future work; the current lumped approximation introduces a systematic bias in the relative key rate comparison that cannot be quantified without the component-level data. The size of this bias is not negligible: Figure 9 shows node loss alone accounting for a 36.7% key-rate reduction for BB84 versus 60.5% for MDI at 25 km, the single largest per-arm contributor to the BB84/MDI gap identified in this work, so a protocol-specific breakdown of  $\Delta_n$  would plausibly shift this margin further rather than narrow it, given MDI’s additional relay-side components.

*NetSquid discrete-event simulation.* The simulations in this work model each photon trial as an independent, identically distributed (i.i.d.) draw: no channel memory, no inter-trial correlations, no time-varying network state. Under this assumption, the key rate distributions produced by NetSquid are in principle reproducible by a simple analytical probability model operating on the same physical parameters. NetSquid’s overhead — launching a DES engine, propagating quantum state objects, managing event queues — is not necessary for the i.i.d. regime. The justification for using NetSquid is extensibility: the DES framework supports temporally correlated dynamics (detector deadtime, jitter, burst traffic, entanglement distillation protocols) that an i.i.d. model cannot accommodate, and these effects become important if this infrastructure is extended to quantum repeater chains or adaptive network protocols. The computational cost of the DES approach is non-negligible: network sweeps across multiple seeds and user counts required on the order of hours of wall-clock time on a multi-core workstation. This constrained the number of seeds and Monte Carlo runs feasible within the project timeline, and limits the statistical confidence of the topology-averaged results. Table 4 summarises representative wall-clock costs on a multi-core workstation: a full point-to-point parameter sweep completes in minutes, but cost grows sharply with network size, reaching hours for a  $N=20$  relay sweep and an estimated weeks for  $N=100$  — infeasible within the project timeline. The dominant cost drivers are the  $\mathcal{O}(N^2)$  user-pair count, the sequential causal event loop within each NetSquid run (which cannot be parallelised across a single trial), and multiprocessing spawn overhead of approximately 2–3 s per Monte Carlo batch, independent of workload.

| Scenario           | Wall-clock time | Feasibility |
|--------------------|-----------------|-------------|
| P2P sweep          | $\sim$ minutes  | OK          |
| $N=20$ relay sweep | $\sim$ hours    | Marginal    |
| $N=100$ sweep      | $\sim$ weeks    | Infeasible  |

Table 4: Wall-clock time for representative NetSquid sweep scenarios on a multi-core workstation.

*Mitigation strategies.* Two mitigations were investigated. Adaptive Monte Carlo sampling terminates trials early once the running key-rate estimate has converged, reserving the full sample budget for slower-converging points; this reduced BB84 wall-clock time by 30% (28% fewer runs,  $<1\%$  relative error) but gave negligible savings for MDI, whose structurally higher variance rarely triggers early stopping. Gaussian process surrogate modelling instead trains a regression model on a coarse NetSquid-generated grid and evaluates a dense  $10^4$ -point grid in under 100 ms, giving near-instant interpolation once trained; however, it requires a substantial upfront training sweep, its accuracy degrades outside the training range, and it defers rather than reduces the underlying simulation cost. Neither approach addresses the root cause, since cost is set by control flow and process management (the sequential event loop and spawn overhead) rather than floating-point throughput — for the same reason, GPU acceleration was considered and rejected, as BB84/MDI operate on 1–2-qubit density matrices ( $2\times 2$  or  $4\times 4$ ) too small for GPU parallelism to outweigh its overhead, and NetSquid has no GPU backend regardless. These constraints motivated the relay-count and user-count ranges explored in Section 3.1.4 and raise a broader question: whether NetSquid’s general-purpose discrete-event engine is the right abstraction for network-scale QKD simulation, or whether quantum causal structure demands bespoke simulation models.

*Metropolitan network scope.* The core BB84-vs-MDI comparison models networks within a  $10 \times 10$  km metropolitan grid, where inter-node distances are short enough for both protocols to establish keys at useful rates without range-extending mechanisms. At wider-area scales — city-to-city or national-backbone links — direct and MDI-QKD links become loss-limited and two distinct range-extending architectures become relevant. Trusted-node relay networks forward classical XOR-combined keys across intermediate nodes; this restores range at the cost of end-to-end security, since a compromised relay node exposes the key (Section 5.1). Quantum repeaters, by contrast, extend range through entanglement swapping and distillation, maintaining end-to-end quantum security without trusted intermediaries, but remain a longer-term technology requiring quantum memories

and mid-circuit measurements at commercially unproven fidelity and scale; this work does not model repeaters. Trusted-node relaying *is* implemented (TBB84, Section 3.1.4) but is exercised only in the real-topology experiments, where range extension is structurally necessary, and deliberately excluded from the core synthetic-grid comparison for the reasons given in Section 5.1. The metropolitan scope of the main comparison isolates the BB84-vs-MDI trade-off from the additional variables introduced by relay-trust forwarding or repeater chain design.

## 5.4 Implications for Protocol Selection

The results support a protocol-selection framework based on two axes: security requirement and scale. At small  $N$  (fewer than approximately  $[N^*]$  users), BB84’s direct mesh is both cheaper and faster than MDI. As  $N$  grows, BB84 fibre cost scales as  $\mathcal{O}(N^2)$  whilst MDI scales as  $\mathcal{O}(N)$ ; MDI becomes cost-competitive at  $[N^*]$  users, a crossover that depends on area, relay count, and component prices. Where the deployment scale exceeds this crossover, the choice between BB84 and MDI reduces to a security question: MDI eliminates detector side-channel vulnerabilities at the cost of a lower key rate; BB84 delivers higher throughput within the trust assumptions of its detector model.

## 6 Conclusion

### 6.1 Summary of Contributions

This work has built a complete simulation and cost modelling infrastructure for the comparative analysis of BB84 and MDI-QKD at metropolitan scale. On the simulation side, both protocols are implemented within the NetSquid discrete-event simulation engine under a ten-layer hardware model that progressively introduces physical imperfections from an idealised baseline to a fully realistic system. Point-to-point key rates are obtained via Monte Carlo simulation across a parameter space covering fibre loss, detector efficiency, dark count rate, TX insertion loss, receiver node loss, source error rate, fibre dephasing, detector basis bias, and BSM beam-splitter efficiency. The simulation infrastructure parallelises across CPU cores using a spawn-context pool and persists all results to a structured SQLite database for post-hoc analysis without re-simulation.

At the network level, both protocols are implemented as multi-user simulators: a direct BB84 mesh and an MDI-QKD network with relay BSM stations. Relay positions are determined by  $k$ -means clustering of random user placements, and all network experiments are repeated across multiple independent seeds to characterise topology variability. A symbolic cost model expresses total CapEx in terms of free parameters  $c_s$  (photon source),  $c_d$  (single-photon detector), and  $c_f$  (fibre per km), with component counts determined by topology. Since manufacturer prices are not publicly available, the model focuses on the structural scaling difference— $\mathcal{O}(N^2)$  fibre for BB84 vs  $\mathcal{O}(N)$  for MDI—and derives the crossover user count  $N^*$  analytically. A stochastic fibre tortuosity model accounts for cable routing overhead in the simulated key rate.

### 6.2 Findings

*[This section will be completed once simulation results are generated. The following bullet points identify the expected finding structure.]*

- The MDI-QKD key rate is approximately  $[X] \times$  lower than BB84 in the bipartite setting at metropolitan distances under layer-9 hardware, consistent with the squaring of channel transmittance imposed by the BSM coincidence requirement.
- The rate gap increases in the network setting due to cross-cluster relay routing detours; the magnitude of the increase depends on relay count and user distribution.

- The symbolic cost model predicts that MDI becomes cost-competitive above a crossover user count  $N^*$ , driven by the  $\mathcal{O}(N)$  vs  $\mathcal{O}(N^2)$  fibre scaling difference; the value of  $N^*$  depends on the ratio  $c_d/c_f$ , relay count  $K$ , and network geometry.
- Detector efficiency is the dominant hardware parameter limiting key rate in the bipartite setting at short and medium distances; fibre attenuation dominates at long distances.

### 6.3 Future Work

Several extensions are of immediate interest.

*Finite-key corrections.* All key rates are asymptotic; finite-key bounds [43] would reduce reported rates and alter the comparison, particularly at long distances where detection events are sparse and block lengths are short.

*Weak coherent pulse sources.* Replacing ideal QD sources with WCP laser sources modelled via the decoy-state technique [27] would give rates more directly comparable to published experimental deployments [11, 16, 17] and would substantially reduce assumed source cost.

*Optimal relay placement.*  $k$ -means relay placement is heuristic. ILP-based placement [22] solves the joint relay placement and routing problem to global optimality, providing tighter lower bounds on minimum MDI deployment cost.

*Brownfield relay expansion.* The growth-model experiments of Section 3.3 and Section B.4 hold the relay count  $K$  fixed while users grow, modelling a provider serving unplanned demand on existing infrastructure. The relay count sweep of Section 3.3, by contrast, still re-clusters and repositions all  $K$  relays from scratch at each  $K$ , which models greenfield network design rather than expanding a live network: a real operator adding capacity has already sunk the cost of its existing  $K$  relays and cannot costlessly relocate them. Extending  $K \rightarrow K + 1$  realistically would instead fix the existing relay positions and solve for where a single new relay should go — and which existing users it should take over — to minimise the resulting total fibre length. This is a harder problem than the placement methods used elsewhere in this work: the new relay’s optimal position depends on which users it serves, and that assignment depends on its position, so it needs its own alternating position/assignment loop (as in `optimise_relays`) restricted to a single new facility with the other  $K$  held fixed, rather than the one-shot strategies used in Section B.4. Not attempted here for want of time.

*Real network topologies.* Extending the simulator to accept fixed network graphs from real city infrastructure would enable direct comparison with published deployed net-

works [11, 16, 17] and provide a practical planning tool for network operators.

*Twin-Field QKD.* TF-QKD and its variants [30] achieve key rates scaling as  $\mathcal{O}(\sqrt{\eta})$ , surpassing the PLOB bound and extending the distance at which MDI-level security is achievable without quantum repeaters. Extending this framework to TF-QKD would complete the repeater-less QKD comparison.

*Operational expenditure.* Incorporating cryogenic OpEx for SNSPD-equipped deployments, using the energetic analysis methodology of [23], would enable full lifetime cost comparisons and may substantially alter the economic case for MDI at high relay counts.

*Trusted-node relay networks.* At wide-area scales where metropolitan link distances are exceeded and direct MDI-QKD links become loss-limited, trusted-node relay architectures become relevant. Extending the cost model and simulation to trusted-node BB84 would enable comparison across the full distance regime.

*Traffic modelling.* A model of key request rates, buffer occupancy, and WDM channel sharing would give a more realistic picture of BB84 and MDI at large  $N$ , where the direct-mesh assumption becomes economically unphysical.



## A Additional Point-to-Point Parameter Sweeps

Section 4.1 discussed four of the eleven point-to-point parameter sweeps produced for this work: distance, detector efficiency, receiver loss, and Charlie’s relay position. Figure 11 reproduces the remaining seven; the discussion below repeatedly compares these against detector efficiency (Figure 5) and receiver loss (Figure 6) from Section 4.1. In every panel, only the swept parameter deviates from its layer-9 default in Table 2; all other parameters are held at that default operating point (20 km Alice–Bob separation, 0.18 dB/km fibre attenuation,  $\eta_d = 0.9$  detector efficiency, 50 cps dark count rate,  $L_i = 0.1$  TX insertion loss, 2.0 dB node loss,  $\varepsilon_s = 0.015$  source error rate,  $\beta = 3.2 \times 10^{-7}$ /km fibre dephasing,  $\eta_{bs} = 0.97$  beam-splitter efficiency, Charlie at the arm midpoint), with 1000 Monte Carlo runs per point and shaded bands showing the sample spread.




Figure 11: Remaining point-to-point parameter sweeps. BB84 in blue (circles), MDI in orange (squares); shaded bands show Monte Carlo sample spread. Except for the parameter swept in each panel, all hardware parameters take their layer-9 default values from Table 2.

Three groupings emerge from these seven sweeps, extending the channel-versus-per-arm pattern established in Section 4.1.

**Channel-side loss.** Fibre attenuation (Figure 11a) behaves exactly like the distance sweep of Section 4.1: the BB84/MDI ratio stays close to  $5.2\text{--}5.3\times$  across the full 0–0.3 dB/km range, confirming that fibre attenuation — like distance — acts on the shared

channel before Alice’s and Bob’s photons diverge towards Charlie, so it scales both protocols together rather than widening the gap between them.

**Per-arm hardware loss.** Beam-splitter efficiency (Figure 11b) and TX insertion loss (Figure 11c) both act after the split, and both widen the ratio, consistent with detector efficiency and receiver loss in Section 4.1. Beam-splitter efficiency is unique among all eleven sweeps in that BB84 is structurally unaffected by it (the dashed reference line in Figure 11b is exactly flat, since  $\eta_{bs}$  only enters MDI’s Bell-state measurement); MDI key rate alone declines from 66.1 kbps at  $\eta_{bs} = 1.0$  to 42.3 kbps at  $\eta_{bs} = 0.8$ , a 36% reduction from a single power of  $\eta_{bs}$  in the coincidence probability. TX insertion loss widens the BB84/MDI ratio from  $4.8\times$  at  $L_i = 0$  to  $9.7\times$  at  $L_i = 0.5$ , the same per-arm doubling pattern seen under detector efficiency (Figure 5) and receiver loss (Section 4.1). X-basis detector efficiency, or basis bias (Figure 11d), belongs to this group too but widens the ratio only modestly, from  $5.0\times$  at  $\eta_X = 1.0$  to  $6.1\times$  at  $\eta_X = 0.5$ , far less than the full detector-efficiency sweep of Figure 5, because here only one of the two measurement bases is degraded rather than both.

**Noise-like parameters.** Dark count rate, fibre dephasing rate, and source error rate (Figures 11e–11g) leave both protocols essentially flat across their full swept ranges at this metropolitan operating point, with the BB84/MDI ratio staying within a few percent of  $5.2\times$  throughout each. None of the three introduces a directional loss term on either arm, so none is expected to shift the ratio; as noted in Section 5.3, this flatness is specific to the signal rates achieved at 20 km and should not be extrapolated to longer, detection-starved links.

## B Relay Placement Optimisation

Section 3.1.4 proposed replacing the plain  $k$ -means-centroid relay placement used elsewhere in this work’s development with a generalised Weiszfeld iteration, on the grounds that it correctly minimises total fibre length — the network’s actual linear-cost objective — rather than the sum of squared distances that  $k$ -means centroids minimise. This appendix works through that argument in more detail: it derives the coupled update of Equation 7 from first principles, works a small example to build intuition for how and why the two objectives diverge, and reports a self-contained geometric comparison of the two placement methods across synthetic topologies matching those used in Section 3.1.4. This comparison is purely geometric — total fibre length under each placement, independent of NetSquid — since, as noted in Section 3.1.4, no network-level key rate results have yet been generated under the new placement method; the discussion questions at the end of this appendix are correspondingly open rather than resolved.

### B.1 Derivation of the Coupled Update

Recall the total fibre length objective from Equation 6,

$$L(x_1, \dots, x_K) = \sum_{k=1}^K \sum_{i \in C_k} \|x_k - p_i\| + \sum_{k < l} \|x_k - x_l\|,$$

where  $C_k$  is the fixed set of users assigned to relay  $k$  and  $x_k$  its free position. Away from the (measure-zero) configurations where  $x_k$  coincides with a user  $p_i$  or another relay  $x_l$ ,  $L$  is differentiable, and each term contributes a unit vector pointing from the fixed point towards  $x_k$ , scaled by the reciprocal distance:

$$\nabla_{x_k} L = \sum_{i \in C_k} \frac{x_k - p_i}{\|x_k - p_i\|^2} + \sum_{l \neq k} \frac{x_k - x_l}{\|x_k - x_l\|^2}.$$

Setting  $\nabla_{x_k} L = 0$  and separating the  $x_k$ -weighted terms from the fixed-point-weighted terms on each side gives

$$x_k \left( \sum_{i \in C_k} \frac{1}{\|x_k - p_i\|} + \sum_{l \neq k} \frac{1}{\|x_k - x_l\|} \right) = \sum_{i \in C_k} \frac{p_i}{\|x_k - p_i\|} + \sum_{l \neq k} \frac{x_l}{\|x_k - x_l\|},$$

which rearranges directly to the fixed-point update of Equation 7 by evaluating the right-hand side at iteration  $t$  and solving for  $x_k$  at iteration  $t + 1$ . This is a majorize-minimise

step, not a gradient step: each relay is repositioned to the weighted mean of the points pulling on it, with weights set by the *previous* iterate’s distances, and iterating is guaranteed to be non-increasing in  $L$  (Kuhn [42]) even though no step size is chosen explicitly.

The degenerate case, where an iterate  $x_k^{(t)}$  lands exactly on a  $p_i$  or  $x_i$ , produces a division by zero, since the corresponding term is non-differentiable there (the norm has a subgradient of unit magnitude in any direction rather than a unique gradient). In practice, and in the numerical experiments below, this is handled by flooring every pairwise distance at a small  $\epsilon$  (we use  $\epsilon = 10^{-9}$  km) before taking its reciprocal; this perturbs the objective negligibly at realistic network scales but avoids the singularity entirely. A more principled fix — checking the subgradient optimality condition explicitly at any iterate landing within  $\epsilon$  of a fixed point, as in the standard single-facility Weiszfeld literature — was not needed at the  $N$ ,  $K$  scales explored here, since exact coincidence between a continuous-uniform random draw and another point has probability zero and was never observed in  $30 \times 9 = 270$  seeded topology trials.

## B.2 Worked Example: Centroid vs. Geometric Median

Figure 12 illustrates the divergence between the two objectives on a small hand-constructed instance: eight users in two well-separated groups of four, one of which (top right) contains a single point noticeably further from the other three than they are from each other. The  $k$ -means centroid of the top-right cluster is pulled visibly towards this one distant point, since it enters the sum-of-squares objective quadratically; the geometric median instead settles closer to the tight group of three, since the same point enters the sum-of-*unsquared*-distances objective only linearly. This is the standard robustness property of the median relative to the mean, familiar from one-dimensional statistics, applied here to a two-dimensional facility-location setting. The effect on total fibre length is modest but directionally consistent with every trial in the larger comparison below: 15.53 km under the  $k$ -means centroid versus 14.62 km under the coupled Weiszfeld update, a 5.9% reduction, obtained without changing which users are assigned to which relay (cluster assignment is identical between the two panels).




Figure 12: A small eight-user, two-relay instance comparing  $k$ -means centroid placement (left) against the generalised Weiszfeld placement (right). Coloured points are users, coloured lines are spoke links, the dashed grey line is the backbone link, and black stars mark relay positions. Cluster assignment is identical in both panels; only relay position differs.

### B.3 Geometric Comparison Across Network Topologies

To quantify the fibre-length reduction at realistic network scales, both placement methods were run on synthetic topologies generated identically to those used in Section 3.1.4:  $N$  users placed uniformly at random within a  $10 \times 10$  km service area, for  $N \in \{10, 20, 50\}$  and relay counts  $K \in \{2, 3, 5\}$ , over 30 independent seeds per  $(N, K)$  pair. For each seed,  $k$ -means clustering ( $n_{\text{init}} = 10$ ) provides both the baseline centroid placement and the initial cluster assignment for the coupled Weiszfeld refinement (Equation 7), which then alternates position updates with nearest-relay reassignment until total fibre length changes by less than  $10^{-9}$  km between iterations (typically several dozen iterations; up to 146 in the smallest, most relay-dense case). Table 5 reports total fibre length — spoke plus backbone, Equation 6 — under each method, averaged across seeds.

| $N$ | $K$ | $k$ -means length (km) | Weiszfeld length (km) | Reduction       |
|-----|-----|------------------------|-----------------------|-----------------|
| 10  | 2   | 29.95                  | 28.57                 | 4.6% $\pm$ 2.7  |
| 10  | 3   | 35.42                  | 30.92                 | 12.6% $\pm$ 3.8 |
| 10  | 5   | 69.58                  | 37.89                 | 45.6% $\pm$ 3.7 |
| 20  | 2   | 57.23                  | 56.00                 | 2.1% $\pm$ 1.4  |
| 20  | 3   | 56.31                  | 53.61                 | 4.8% $\pm$ 1.8  |
| 20  | 5   | 84.80                  | 67.13                 | 20.8% $\pm$ 4.6 |
| 50  | 2   | 143.01                 | 142.11                | 0.6% $\pm$ 0.4  |
| 50  | 3   | 124.03                 | 122.06                | 1.6% $\pm$ 0.8  |
| 50  | 5   | 133.63                 | 126.01                | 5.7% $\pm$ 1.9  |

Table 5: Mean total fibre length (spoke + backbone) under  $k$ -means centroid versus generalised Weiszfeld relay placement, across 30 seeds per  $(N, K)$  pair,  $10 \times 10$  km service area. Reduction is  $100(L_{\text{kmeans}} - L_{\text{Weiszfeld}})/L_{\text{kmeans}}$ , mean  $\pm$  one standard deviation across seeds.

The reduction is positive in every trial — consistent with the monotonic non-increase guarantee noted above, since the Weiszfeld iteration starts from the  $k$ -means solution and never leaves it worse off — and grows sharply with  $K$  at fixed  $N$ : from 4.6% at  $K = 2$  to 45.6% at  $K = 5$  for  $N = 10$ , and from 0.6% to 5.7% across the same  $K$  range for  $N = 50$ . This follows directly from Equation 6: the backbone term has  $K(K - 1)/2$  pairwise contributions, growing quadratically in  $K$ , while the spoke term grows only linearly in  $N$ ; at larger  $K$  the backbone increasingly dominates total cost, and it is precisely the backbone coupling that a per-cluster centroid ignores entirely. Conversely, the reduction shrinks as  $N$  grows at fixed  $K$  — from 45.6% to 5.7% at  $K = 5$  between  $N = 10$  and  $N = 50$  — because a larger user base dilutes each relay’s positioning error relative to the now spoke-dominated total, and because tighter, more populated clusters leave less room for a single geometric outlier to pull the centroid away from the coupled optimum. The practical implication is that this refinement matters most exactly where the MDI network topology already carries a real relay count  $K$  relative to  $N$  — the regime the cost-crossover analysis of Section 4.3 identifies as commercially relevant for smaller deployments before BB84’s  $\mathcal{O}(N^2)$  fibre term dominates — and least where  $N$  is large enough that per-user fibre cost already dominates the network total regardless of relay placement.

## B.4 Network-Level Comparison Under Organic Growth

The geometric comparison above holds cluster assignment and  $N$  fixed. A second, network-level comparison (`relay_placement.py`) instead asks how different placement strategies compare as a network grows under the fixed-membership provider-growth model of Section 3.3:  $N_{\min}$  users are drawn into  $K$  fixed catchments, each relay is placed once by the strategy under test, and the relay position is then *frozen* while further users join one at a time, each assigned to a uniformly-random existing catchment. Three strategies are compared: *centroid*, the mean position of a catchment’s users; *boundary*, the centroid displaced by one standard deviation of the catchment’s own spread, in the direction of the population-weighted mean of the other catchments’ centroids (a heuristic proxy for the backbone term, without solving for it); and *Weiszfeld*, Equation 7 with catchment membership held fixed rather than reassigned. Experiment 1 runs this comparison at  $K = 1$  (a single catchment spanning the whole service area), where only centroid and Weiszfeld apply — boundary is undefined without a second catchment to displace toward, and Weiszfeld reduces to the plain single-facility geometric median of Equation 5 since the backbone term in Equation 7 vanishes at  $K = 1$ . Experiment 2 runs it at  $K = 2$ , comparing all three strategies. Because the growth sequence, catchment split, and catchment membership are identical across strategies for a given seed — only where each strategy chooses to place its relay at  $N_{\min}$  differs — the comparison isolates the effect of the initial placement decision on robustness to unplanned subsequent growth, rather than re-answering the already-resolved question of which placement is instantaneously optimal for a fixed user set.

## B.5 Open Questions

The comparison above resolves the geometric question — which placement minimises total fibre length for a *fixed* backbone topology and cluster assignment — but leaves several modelling questions open, appropriate for discussion before the new placement method is adopted for the network-level key rate sweeps of Section 3.1.4:

- *Backbone topology.* Table 3 assumes a full  $K(K - 1)/2$  relay mesh. A minimum spanning tree or Steiner-tree backbone would in general have strictly lower total length for  $K > 2$ , and would change the placement problem itself (the coupling term in Equation 6 would run over tree edges rather than all pairs). Whether the full mesh is retained for redundancy/latency reasons independent of cost, or replaced with a cheaper backbone topology, should be decided before the cost model and placement method are finalised together.

- *Assignment optimality.* Cluster reassignment after each position update is by nearest Euclidean distance to a relay, which is the correct assignment rule for the spoke term alone but does not account for the backbone term: a user could in principle be assigned to a farther relay if doing so allowed both relays to move to a jointly cheaper configuration. This is a combinatorial partitioning question distinct from the continuous position optimisation solved by Equation 7, and was not explored here.
- *Convergence cost.* The Weiszfeld refinement took up to 146 outer iterations to reach a  $10^{-9}$  km convergence tolerance at the smallest, most relay-dense configuration tested ( $N = 10$ ,  $K = 5$ ); this is negligible against NetSquid’s own per-trial cost (Section 5.3) but would be worth profiling at the  $N = 100$  scale flagged there as currently infeasible.
- *Net effect on key rate.* A shorter mean path length should translate into a modest key rate improvement through reduced fibre attenuation (Section 3.1.2, layer 1), but by how much depends on where in the loss-versus-distance curve of Figure 8 the affected links sit; this can only be answered by re-running the network-level sweeps of Section 3.1.4 under the new placement, which had not yet been done at the time of writing.
- *Degenerate coincidence.* Both Equations 5 and 7 divide by the distance between the current iterate and each fixed point, which is undefined if a relay iterate ever coincides exactly with a user (or, in the multi-facility case, another relay). The implementation guards against this by flooring the distance at  $10^{-9}$  km rather than applying Ostresh’s modification — the standard fix, which tests whether the coincident point is itself already a stationary point of the weighted sum and, if not, steps away from it along a descent direction. The floor avoids a division-by-zero crash but is not equivalent: it lets that term’s weight grow to  $\mathcal{O}(10^9)$ , which can pin the next iterate back near the same point rather than resolving the singularity correctly. Coincidence is measure-zero under the continuous uniform and Gaussian user placements used throughout this work, so it was not observed to trigger in practice, but the implementation should not be trusted as-is on discretised or clustered location data (e.g. users snapped to a road network, or a relay initialised exactly at a fixed real-topology co-ordinate) where exact coincidence is plausible.
- *Which topologies favour Weiszfeld.* Weiszfeld does not uniformly beat centroid or boundary placement at the network-key-rate level: across the growth-model com-



parison of Section B.4, its ranking against the other two strategies varies from seed to seed and from  $N$  to  $N$ , sometimes best, sometimes not. This is expected in principle — centroid minimises spoke cost per catchment in isolation, while Weiszfeld additionally couples in the backbone term, so the two objectives should diverge most when the backbone is a non-negligible fraction of total fibre length (catchments close together relative to their spread, or markedly unequal in size) and coincide when it is not. Since catchment membership for  $N > N_{\min}$  is now assigned by an independent uniform coin flip per new user rather than a fixed even split, catchment size imbalance is a first-class, randomly-varying outcome of the growth process itself rather than a fixed property of the initial topology — strengthening the case that it is a real driver of the mixed ranking, though this work only observed the mixed outcome without identifying which geometric property of a given seed predicts it. A natural follow-up would be to regress the per-seed key rate gap between strategies against candidate covariates (cluster size imbalance realised by  $N_{\max}$ , inter-cluster separation relative to intra-cluster spread, backbone-to-spoke cost ratio at the initial placement) using the per-seed data already logged by `relay_placement.py`, and, if a clear predictor emerges, confirm it with a dedicated sweep that varies that one covariate while holding the others fixed. Not undertaken here for want of time.

## References

- [1] R. L. Rivest, A. Shamir, and L. Adleman, “A method for obtaining digital signatures and public-key cryptosystems,” *Commun. ACM*, vol. 21, p. 120–126, Feb. 1978.
- [2] N. Koblitz, “Elliptic Curve Cryptosystems,” *Mathematics of Computation*, vol. 48, pp. 203–209, Jan. 1987.
- [3] V. S. Miller, “Use of Elliptic Curves in Cryptography,” in *Advances in Cryptology — CRYPTO ’85 Proceedings* (H. C. Williams, ed.), (Berlin, Heidelberg), pp. 417–426, Springer Berlin Heidelberg, 1986.
- [4] P. W. Shor, “Polynomial-Time Algorithms for Prime Factorization and Discrete Logarithms on a Quantum Computer,” *SIAM Journal on Computing*, vol. 26, p. 1484–1509, Oct. 1997.
- [5] S. Wiesner, “Conjugate coding,” *SIGACT News*, vol. 15, p. 78–88, Jan. 1983.
- [6] C. H. Bennett and G. Brassard, “Quantum cryptography: Public key distribution and coin tossing,” *Theoretical Computer Science*, vol. 560, p. 7–11, Dec. 2014.
- [7] Y.-A. Chen *et al.*, “An integrated space-to-ground quantum communication network over 4,600 kilometres,” *Nature*, vol. 589, 01 2021.
- [8] M. Peev *et al.*, “The SECOQC quantum key distribution network in Vienna,” *New Journal of Physics*, vol. 11, p. 075001, 07 2009.
- [9] M. Sasaki *et al.*, “Field test of quantum key distribution in the Tokyo QKD Network,” *Optics Express*, vol. 19, p. 10387, May 2011.
- [10] H.-K. Lo, M. Curty, and B. Qi, “Measurement-Device-Independent Quantum Key Distribution,” *Phys. Rev. Lett.*, vol. 108, p. 130503, Mar 2012.
- [11] J. Dynes *et al.*, “Cambridge quantum network,” *npj Quantum Information*, vol. 5, 12 2019.
- [12] W. Yan, X. Zheng, W. Wen, L. Lu, Y. Du, Y. Lu, S. Zhu, and X.-S. Ma, “A measurement-device-independent quantum key distribution network using optical frequency comb,” 2025.
- [13] C. Elliott, “Building the quantum network,” *New Journal of Physics*, vol. 4, p. 46, July 2002.

- [14] H. Qin *et al.*, “The National Quantum-Safe Network in Singapore,” in *Proceedings of QCrypt 2024*, 2024.
- [15] M. Pittaluga, Y. S. Lo, A. Brzosko, R. I. Woodward, M. S. Winnel, T. Roger, J. F. Dynes, K. A. Owen, S. Juarez, P. Rydlichowski, D. Vicinanza, G. Roberts, and A. J. Shields, “Coherent Quantum Communications Across National Scale Telecommunication Infrastructure,” 2024.
- [16] Y.-L. Tang *et al.*, “Measurement-Device-Independent Quantum Key Distribution over Untrustful Metropolitan Network,” *Physical Review X*, vol. 6, Mar. 2016.
- [17] R. C. Berrevoets *et al.*, “Deployed measurement-device independent quantum key distribution and Bell-state measurements coexisting with standard internet data and networking equipment,” *Communications Physics*, vol. 5, July 2022.
- [18] Q\*Bird, “Fortifying Europe’s quantum communication – Q\*Bird and BeQCI achieve the first cross-border MDI QKD link in Benelux,” 2025.
- [19] V. Martin *et al.*, “MadQCI: a heterogeneous and scalable SDN-QKD network deployed in production facilities,” *npj Quantum Information*, vol. 10, 09 2024.
- [20] E. Commission, “Open European Quantum Key Distribution Testbed,” 2024.
- [21] R. Yehia, S. Neves, E. Diamanti, and I. Kerenidis, “Quantum City: simulation of a practical near-term metropolitan quantum network,” 2022.
- [22] V. Karavias, E. Hugues-Salas, A. Beghelli, A. Lord, and M. C. Payne, “Comparative Cost Analysis of Prepare and Measure and Entanglement-Based QKD Networks,” in *International Conference on Optical Network Design and Modeling, ONDM, 2025*, 2025.
- [23] R. Yehia, Y. Piétri, C. Pascual-García, P. Lefebvre, and F. Centrone, “Energetic Analysis of Emerging Quantum Communication Protocols,” 2025.
- [24] T. Coopmans, R. Knegjens, A. Dahlberg, D. Maier, L. Nijsten, J. de Oliveira Filho, M. Papendrecht, J. Rabbie, F. Rozpędek, M. Skrzypczyk, L. Wubben, W. de Jong, D. Podareanu, A. Torres-Knoop, D. Elkouss, and S. Wehner, “NetSquid, a NETwork Simulator for QUantum Information using Discrete events,” *Communications Physics*, vol. 4, no. 1, 2021.
- [25] P. W. Shor and J. Preskill, “Simple Proof of Security of the BB84 Quantum Key Distribution Protocol,” *Physical Review Letters*, vol. 85, p. 441–444, July 2000.

- [26] E. I. T. C. Institute, “What is the Photon Number Splitting (PNS) attack, and how does it constrain the communication distance in quantum cryptography?,” 2024.
- [27] H.-K. Lo, X. Ma, and K. Chen, “Decoy State Quantum Key Distribution,” *Phys. Rev. Lett.*, vol. 94, p. 230504, Jun 2005.
- [28] C. K. Hong, Z. Y. Ou, and L. Mandel, “Measurement of subpicosecond time intervals between two photons by interference,” *Phys. Rev. Lett.*, vol. 59, pp. 2044–2046, Nov. 1987.
- [29] S. Pirandola, R. Laurenza, C. Ottaviani, and L. Banchi, “Fundamental limits of repeaterless quantum communications,” *Nature Communications*, vol. 8, Apr. 2017.
- [30] M. Lucamarini, Z. L. Yuan, J. F. Dynes, and A. J. Shields, “Overcoming the rate–distance limit of quantum key distribution without quantum repeaters,” *Nature*, vol. 557, p. 400–403, May 2018.
- [31] Corning, “SMF-28 Ultra Optical Fiber – Product Information Sheet PI-1424-AEN.” Product datasheet, 2023.
- [32] ID Quantique, “ID281 Superconducting Nanowire Single-Photon Detection System – Product Datasheet.” Product datasheet, 2024.
- [33] OZ Optics, “Polarization Maintaining Fused Fiber Couplers/Splitters – Product Datasheet (DTS0092).” Product datasheet, 2013.
- [34] A. Duplinskiy, V. Ustimchik, A. Kanapin, V. Kurochkin, and Y. Kurochkin, “Low loss QKD optical scheme for fast polarization encoding,” *Optics Express*, vol. 25, no. 23, pp. 28886–28897, 2017.
- [35] Quandela, “Prometheus Single-Photon Source – Product Datasheet.” Product datasheet, 2024.
- [36] F. Grasselli, G. Chesi, N. Walk, H. Kampermann, A. Widomski, M. Ogrodnik, M. Karpiński, C. Macchiavello, D. Bruß, and N. Wyderka, “Quantum Key Distribution with Basis-Dependent Detection Probability,” *Physical Review Applied*, vol. 23, p. 044011, 2025.
- [37] Thorlabs, “TN1550R5F2 – Narrowband Fiber Optic Coupler, 1550 nm, 50:50 Ratio – Product Datasheet (TTN168601-S01, Rev C).” Product datasheet, 2021.

- [38] ID Quantique, “ID230 Infrared Single-Photon Detector – Product Datasheet.” Product datasheet, 2024.
- [39] W. Miehle, “Link-Length Minimization in Networks,” *Operations Research*, vol. 6, no. 2, pp. 232–243, 1958.
- [40] M. E. O’Kelly, “The Location of Interacting Hub Facilities,” *Transportation Science*, vol. 20, no. 2, pp. 92–106, 1986.
- [41] S. Alumur and B. Y. Kara, “Network hub location problems: The state of the art,” *European Journal of Operational Research*, vol. 190, no. 1, pp. 1–21, 2008.
- [42] H. W. Kuhn, “A note on Fermat’s problem,” *Mathematical Programming*, vol. 4, no. 1, pp. 98–107, 1973.
- [43] T.-T. Song, Q.-Y. Wen, F.-Z. Guo, and X.-Q. Tan, “Finite-key analysis for measurement-device-independent quantum key distribution,” *Phys. Rev. A*, vol. 86, p. 022332, Aug 2012.